

Table of contents

From neural networks...

The basic ideas

Logistic regression and one layer neural-network

Convolutional neural networks

Graph neural networks

Learning with graphs

What is a GNN ?

A bit of group theory

Invariance and equivariance

Permutation invariance/equivariance

Message-passing neural networks

Examples of GNN

The whole pipeline

Expressivity of GNN

Conclusion

More on unsupervised node embeddings techniques

Unsupervised node embeddings

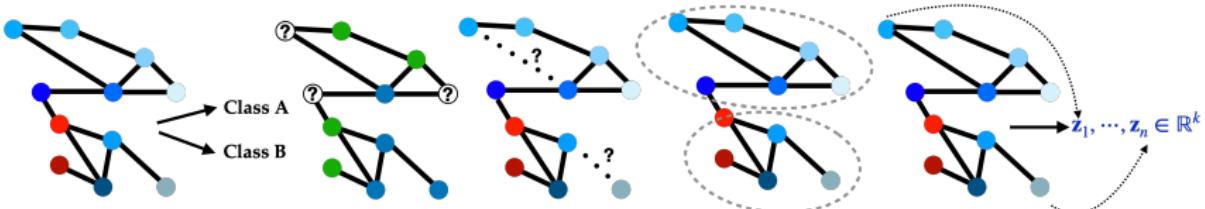
Frameworks considered here

Supervised:

- ▶ Graph classification: labelled graphs → label new graph (molecule classification, drug efficiency prediction).
- ▶ Node (or edge) classification: labelled nodes → label other nodes (advertisement, protein interface prediction).

Unsupervised (semi-supervised):

- ▶ Community detection: one graph → group nodes (social network analysis).
- ▶ Link prediction: one graph → potential new edge.
- ▶ Unsupervised node embeddings.



Some limitations

Tip of the iceberg

- ▶ Approx. 100 GNN papers a month on arXiv.
- ▶ Despite 1000s of papers, same ideas coming round: be critical, learn to spot incremental changes!
- ▶ We will only see the most well-known architectures (according to me).
- ▶ Be aware that it might already be out-of-date.
- ▶ Some surveys [Wu et al. 2021](#); [Zhang, Cui, and Zhu 2020](#); [William L Hamilton 2020](#).
- ▶ See also <https://github.com/houchengbin/awesome-GNN-papers>.

Table of contents

From neural networks...

The basic ideas

Logistic regression and one layer neural-network

Convolutional neural networks

Graph neural networks

Learning with graphs

What is a GNN ?

A bit of group theory

Invariance and equivariance

Permutation invariance/equivariance

Message-passing neural networks

Examples of GNN

The whole pipeline

Expressivity of GNN

Conclusion

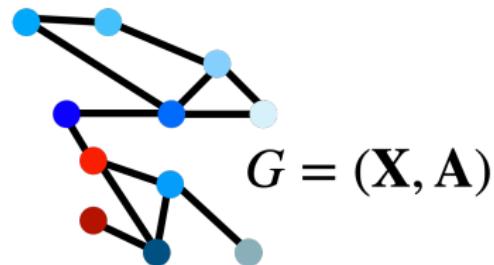
More on unsupervised node embeddings techniques

Unsupervised node embeddings

What is a graph neural network ?

Framework

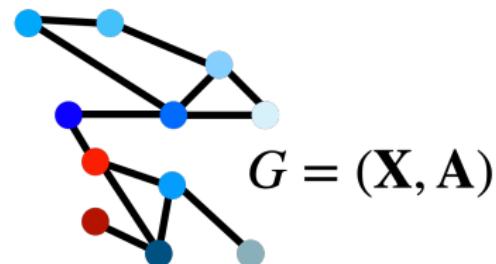
- ▶ Graphs considered here:
- ▶ $G = (V, E)$ with $|V| = n$, features on the nodes.
- ▶ Adjacency matrix $\mathbf{A} \in \mathbb{R}^{n \times n}$.
- ▶ Feature matrix $\mathbf{X} \in \mathbb{R}^{n \times d}$, feature $\mathbf{x}_i \in \mathbb{R}^d$.



What is a graph neural network ?

Framework

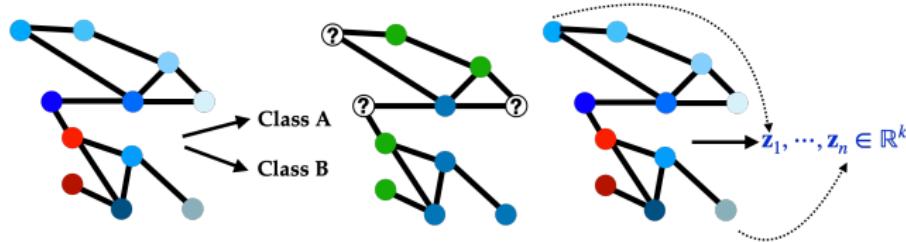
- ▶ Graphs considered here:
- ▶ $G = (V, E)$ with $|V| = n$, features on the nodes.
- ▶ Adjacency matrix $\mathbf{A} \in \mathbb{R}^{n \times n}$.
- ▶ Feature matrix $\mathbf{X} \in \mathbb{R}^{n \times d}$, feature $\mathbf{x}_i \in \mathbb{R}^d$.



What is a GNN ?

A **specific parametrized function** that takes a input a graph $G = (\mathbf{X}, \mathbf{A})$

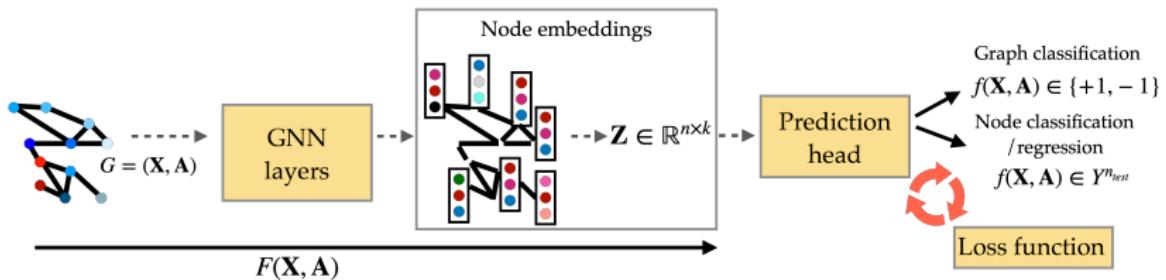
- ▶ It is made of a **combination of different layers**.
- ▶ Graph classification, node classification/regression, node embedding



- ▶ Notations: vector output $f(\mathbf{X}, \mathbf{A})$, matrix output $F(\mathbf{X}, \mathbf{A})$.

What properties to ensure ?

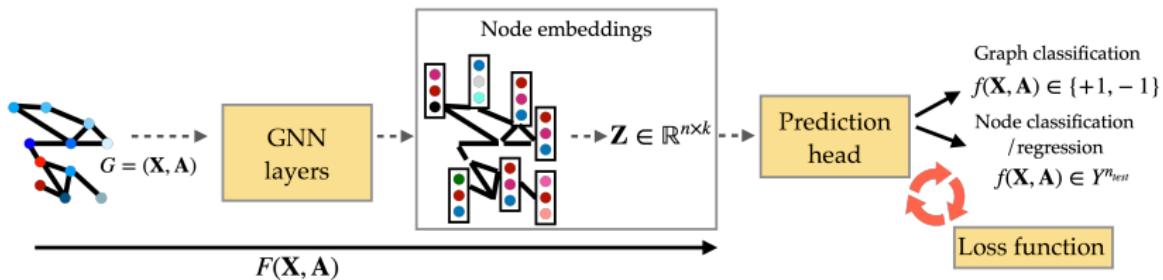
The training pipeline



- ▶ Always same procedure: find an embedding of the nodes $F(\mathbf{X}, \mathbf{A}) \in \mathbb{R}^{n \times k}$ and then do stuff.

What properties to ensure ?

The training pipeline



- ▶ Always same procedure: find an embedding of the nodes $F(\mathbf{X}, \mathbf{A}) \in \mathbb{R}^{n \times k}$ and then do stuff.

Properties to ensure

- ▶ Prediction on the graph level: the output $f(\mathbf{X}, \mathbf{A})$ must be **invariant to permutations of the graph**.
- ▶ Prediction on the node level: permutation of the graph should produce a different result but **be predictable**.
- ▶ It will be formalized with the notion of invariance/equivariance.

Table of contents

From neural networks...

The basic ideas

Logistic regression and one layer neural-network

Convolutional neural networks

Graph neural networks

Learning with graphs

What is a GNN ?

A bit of group theory

Invariance and equivariance

Permutation invariance/equivariance

Message-passing neural networks

Examples of GNN

The whole pipeline

Expressivity of GNN

Conclusion

More on unsupervised node embeddings techniques

Unsupervised node embeddings

On invariance and equivariance

The right symmetries can facilitate learning

- ▶ Fit a polynomial $\hat{f}(x) = \sum_{n=0}^N \theta_n x^n$ on

symmetric: $f(-x) = f(x)$

antisymmetric: $f(-x) = -f(x)$

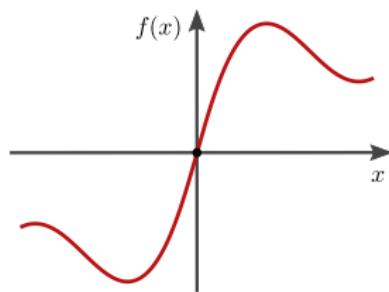
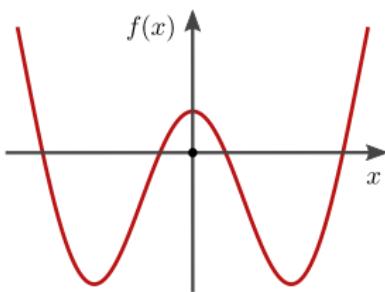


Figure: From Weiler et al. 2023

- ▶ Ignore prior knowledge about the function.
- ▶ Better: fit $\sum_{n \text{ even}}^N \theta_n x^n$ (invariant) or $\sum_{n \text{ odd}}^N \theta_n x^n$ (equivariant).
- ▶ Need half of the parameters + generalize well.

On invariance and equivariance

On the previous episodes

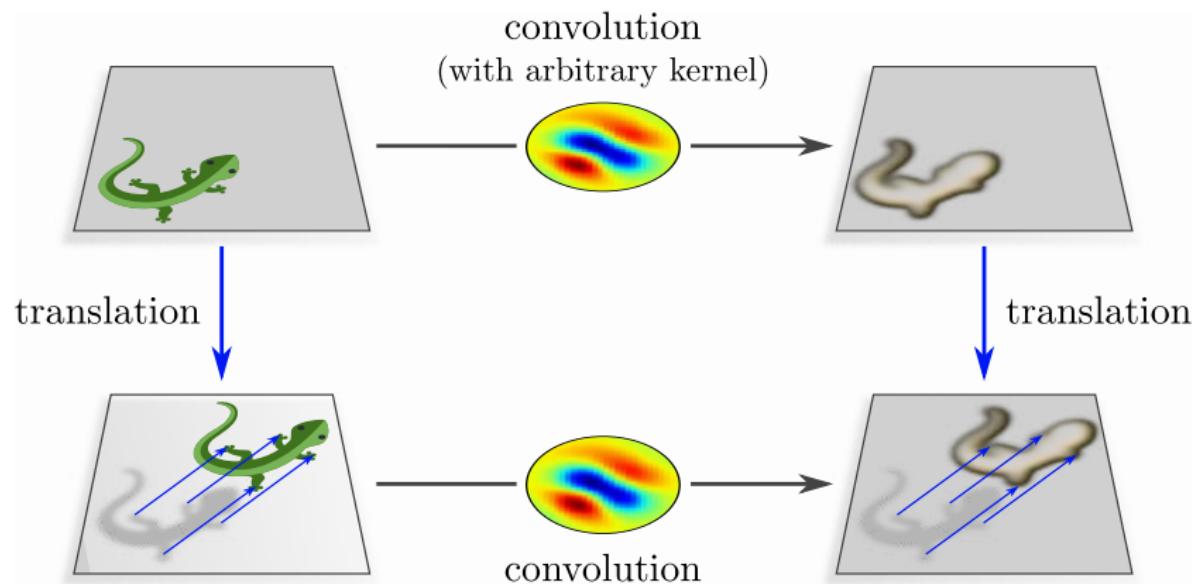


Figure: From Weiler et al. 2023

On invariance and equivariance

A little bit of group theory

A group \mathfrak{G} is a set along with a binary operation $\circ : \mathfrak{G} \times \mathfrak{G} \rightarrow \mathfrak{G}$ satisfying

- ▶ *Associativity*: $\forall g, h, i \in \mathfrak{G}, (g \circ h) \circ i = g \circ (h \circ i)$.
- ▶ *Identity*: there exists $e \in \mathfrak{G}$ such that $\forall g \in \mathfrak{G}, g \circ e = e \circ g = g$.
- ▶ *Inverse*: For each $g \in \mathfrak{G}$ there exists $g^{-1} \in \mathfrak{G}$ such that $g \circ g^{-1} = g^{-1} \circ g = e$.
- ▶ *Closure*: $\forall g, h \in \mathfrak{G}, g \circ h \in \mathfrak{G}$.

Commutativity is not part of this definition ($g \circ h \neq h \circ g$).

On invariance and equivariance

A little bit of group theory

A group \mathfrak{G} is a set along with a binary operation $\circ : \mathfrak{G} \times \mathfrak{G} \rightarrow \mathfrak{G}$ satisfying

- ▶ *Associativity*: $\forall g, h, i \in \mathfrak{G}, (g \circ h) \circ i = g \circ (h \circ i)$.
- ▶ *Identity*: there exists $e \in \mathfrak{G}$ such that $\forall g \in \mathfrak{G}, g \circ e = e \circ g = g$.
- ▶ *Inverse*: For each $g \in \mathfrak{G}$ there exists $g^{-1} \in \mathfrak{G}$ such that $g \circ g^{-1} = g^{-1} \circ g = e$.
- ▶ *Closure*: $\forall g, h \in \mathfrak{G}, g \circ h \in \mathfrak{G}$.

Commutativity is not part of this definition ($g \circ h \neq h \circ g$).

Some examples

- ▶ Translation group on \mathbb{Z}^2 is an Abelian group:

$$(m, n) \circ (p, q) = (n + p, m + q).$$

- ▶ Translation + rotations, mirror reflections.
- ▶ Permutation group $S_n = \{\sigma : \llbracket n \rrbracket \rightarrow \llbracket n \rrbracket, \sigma \text{ is a bijection}\}$ with the composition of functions.

On invariance and equivariance

Group action

Given a set Ω and a group \mathfrak{G} , a (left) **group action** of \mathfrak{G} on Ω is a function

$$\begin{aligned}\mathfrak{G} \times \Omega &\rightarrow \Omega \\ (\mathfrak{g}, x) &\rightarrow \mathfrak{g}x\end{aligned}$$

satisfying

- ▶ $\forall x \in \Omega, \mathfrak{e}x = x$
- ▶ *Compatibility:* $\forall \mathfrak{g}, \mathfrak{h} \in \mathfrak{G}, \forall x \in \Omega, \mathfrak{g}(\mathfrak{h}x) = (\mathfrak{g} \circ \mathfrak{h})x.$
- ▶ It acts on the element of the sets via the group.
- ▶ **A set endowed with an action of \mathfrak{G} on it is called a \mathfrak{G} -set.**

On invariance and equivariance

Group action

Given a set Ω and a group \mathfrak{G} , a (left) **group action** of \mathfrak{G} on Ω is a function

$$\begin{aligned}\mathfrak{G} \times \Omega &\rightarrow \Omega \\ (\mathfrak{g}, x) &\mapsto \mathfrak{g}x\end{aligned}$$

satisfying

- ▶ $\forall x \in \Omega, \mathfrak{e}x = x$
- ▶ *Compatibility:* $\forall \mathfrak{g}, \mathfrak{h} \in \mathfrak{G}, \forall x \in \Omega, \mathfrak{g}(\mathfrak{h}x) = (\mathfrak{g} \circ \mathfrak{h})x.$
- ▶ It acts on the element of the sets via the group.
- ▶ **A set endowed with an action of \mathfrak{G} on it is called a \mathfrak{G} -set.**

Translation of functions

- ▶ Group of translations $\mathfrak{G} = \{\tau_x, x \in \mathbb{R}\}$ with $\tau_x \circ \tau_y = \tau_{x+y}$. Identity element τ_0 .
- ▶ For a function f and τ_x the group action

$$\tau_x f := t \rightarrow f(t - x).$$

On invariance and equivariance

Group action

Given a set Ω and a group \mathfrak{G} , a (left) **group action** of \mathfrak{G} on Ω is a function

$$\mathfrak{G} \times \Omega \rightarrow \Omega$$

$$(\mathfrak{g}, x) \mapsto \mathfrak{g}x$$

satisfying

- ▶ $\forall x \in \Omega, \mathfrak{e}x = x$
- ▶ *Compatibility:* $\forall \mathfrak{g}, \mathfrak{h} \in \mathfrak{G}, \forall x \in \Omega, \mathfrak{g}(\mathfrak{h}x) = (\mathfrak{g} \circ \mathfrak{h})x.$
- ▶ It acts on the element of the sets via the group.
- ▶ **A set endowed with an action of \mathfrak{G} on it is called a \mathfrak{G} -set.**

Permutation of vectors

- ▶ Group of permutations S_n with composition \circ . Identity element id .
- ▶ For $\mathbf{x} \in \mathbb{R}^n$ a group action is $\sigma\mathbf{x} = (x_{\sigma(1)}, x_{\sigma(2)}, \dots, x_{\sigma(n)}).$
- ▶ Is it a left group action ?

On invariance and equivariance

Group action

Given a set Ω and a group \mathfrak{G} , a (left) **group action** of \mathfrak{G} on Ω is a function

$$\begin{aligned}\mathfrak{G} \times \Omega &\rightarrow \Omega \\ (\mathfrak{g}, x) &\mapsto \mathfrak{g}x\end{aligned}$$

satisfying

- ▶ $\forall x \in \Omega, \mathfrak{e}x = x$
- ▶ *Compatibility:* $\forall \mathfrak{g}, \mathfrak{h} \in \mathfrak{G}, \forall x \in \Omega, \mathfrak{g}(\mathfrak{h}x) = (\mathfrak{g} \circ \mathfrak{h})x.$
- ▶ It acts on the element of the sets via the group.
- ▶ **A set endowed with an action of \mathfrak{G} on it is called a \mathfrak{G} -set.**

Permutation of vectors

- ▶ For $\mathbf{x} \in \mathbb{R}^n$ a group action is $\sigma\mathbf{x} = (x_{\sigma(1)}, x_{\sigma(2)}, \dots, x_{\sigma(n)}).$
- ▶ Def $(\sigma_1\mathbf{x})_i = x_{\sigma_1(i)}$. So $(\sigma_2(\sigma_1\mathbf{x}))_i = (\sigma_1\mathbf{x})_{\sigma_2(i)} = x_{\sigma_1(\sigma_2(i))} = x_{\sigma_1 \circ \sigma_2(i)}.$
- ▶ Thus $\sigma_2(\sigma_1\mathbf{x}) = (\sigma_1 \circ \sigma_2)\mathbf{x} \neq (\sigma_2 \circ \sigma_1)\mathbf{x}.$

On invariance and equivariance

Group action

Given a set Ω and a group \mathfrak{G} , a (left) **group action** of \mathfrak{G} on Ω is a function

$$\begin{aligned}\mathfrak{G} \times \Omega &\rightarrow \Omega \\ (\mathfrak{g}, x) &\rightarrow \mathfrak{g}x\end{aligned}$$

satisfying

- ▶ $\forall x \in \Omega, \mathfrak{e}x = x$
- ▶ *Compatibility:* $\forall \mathfrak{g}, \mathfrak{h} \in \mathfrak{G}, \forall x \in \Omega, \mathfrak{g}(\mathfrak{h}x) = (\mathfrak{g} \circ \mathfrak{h})x.$
- ▶ It acts on the element of the sets via the group.
- ▶ **A set endowed with an action of \mathfrak{G} on it is called a \mathfrak{G} -set.**

Permutation of vectors

- ▶ For $\mathbf{x} \in \mathbb{R}^n$ a **left** group action is $\sigma\mathbf{x} = (x_{\sigma^{-1}(1)}, x_{\sigma^{-1}(2)}, \dots, x_{\sigma^{-1}(n)}).$
- ▶ Def $(\sigma_1\mathbf{x})_i = x_{\sigma_1^{-1}(i)}$. So
 $(\sigma_2(\sigma_1\mathbf{x}))_i = (\sigma_1\mathbf{x})_{\sigma_2^{-1}(i)} = x_{\sigma_1^{-1}(\sigma_2^{-1}(i))} = x_{(\sigma_2 \circ \sigma_1)^{-1}(i)}.$
- ▶ Thus $\sigma_2(\sigma_1\mathbf{x}) = (\sigma_2 \circ \sigma_1)\mathbf{x}.$

Table of contents

From neural networks...

The basic ideas

Logistic regression and one layer neural-network

Convolutional neural networks

Graph neural networks

Learning with graphs

What is a GNN ?

A bit of group theory

Invariance and equivariance

Permutation invariance/equivariance

Message-passing neural networks

Examples of GNN

The whole pipeline

Expressivity of GNN

Conclusion

More on unsupervised node embeddings techniques

Unsupervised node embeddings

A formal definition of invariance

Invariance

Let Ω be a \mathfrak{G} -set. A function $f : \Omega \rightarrow Y$ is **\mathfrak{G} -invariant** if

$$\forall x \in \Omega, \forall g \in \mathfrak{G}, f(gx) = f(x).$$

- ▶ f is \mathfrak{G} -invariant if its output is unaffected by the group action.

A formal definition of invariance

Invariance

Let Ω be a \mathfrak{G} -set. A function $f : \Omega \rightarrow Y$ is **\mathfrak{G} -invariant** if

$$\forall x \in \Omega, \forall g \in \mathfrak{G}, f(gx) = f(x).$$

- ▶ f is \mathfrak{G} -invariant if its output is unaffected by the group action.

Permutation invariant functions

Find three functions $f, g, h : \mathbb{R}^n \rightarrow \mathbb{R}$ that are S_n -invariant.

A formal definition of invariance

Invariance

Let Ω be a \mathfrak{G} -set. A function $f : \Omega \rightarrow Y$ is **\mathfrak{G} -invariant** if

$$\forall x \in \Omega, \forall g \in \mathfrak{G}, f(gx) = f(x).$$

- ▶ f is \mathfrak{G} -invariant if its output is unaffected by the group action.

Permutation invariant functions

- ▶ $f(\mathbf{x}) = \sum_{i=1}^n x_i, g(\mathbf{x}) = \max_{i \in [n]} x_i, h(\mathbf{x}) = \text{sort}(\mathbf{x})$ (to \mathbb{R}^n).
- ▶ Characterization of all linear permutation invariant functions
 $L : \mathbb{R}^{n^k} \rightarrow \mathbb{R}$ ([Maron et al. 2018](#)).

A formal definition of invariance

Invariance

Let Ω be a \mathfrak{G} -set. A function $f : \Omega \rightarrow Y$ is **\mathfrak{G} -invariant** if

$$\forall x \in \Omega, \forall g \in \mathfrak{G}, f(gx) = f(x).$$

- ▶ f is \mathfrak{G} -invariant if its output is unaffected by the group action.

Permutation invariant functions

Let $\mathbf{X} \in \mathbb{R}^{n \times d}$. The action of σ on \mathbf{X} is $\sigma\mathbf{X} = (X_{\sigma^{-1}(i)j})_{ij}$. Find a permutation invariant function $F : \mathbb{R}^{n \times d} \rightarrow \mathbb{R}$.

A formal definition of invariance

Invariance

Let Ω be a \mathfrak{G} -set. A function $f : \Omega \rightarrow Y$ is **\mathfrak{G} -invariant** if

$$\forall x \in \Omega, \forall g \in \mathfrak{G}, f(gx) = f(x).$$

- ▶ f is \mathfrak{G} -invariant if its output is unaffected by the group action.

Permutation invariant functions

Let $\mathbf{X} \in \mathbb{R}^{n \times d}$. The action of σ on \mathbf{X} is $\sigma\mathbf{X} = (X_{\sigma^{-1}(i)j})_{ij}$. Find a permutation invariant function $F : \mathbb{R}^{n \times d} \rightarrow \mathbb{R}$.

- ▶ With $\mathbf{X} = (\mathbf{x}_1, \dots, \mathbf{x}_n)^\top$ and $F(\mathbf{X}) = \phi(\sum_{i=1}^n \psi(\mathbf{x}_i))$ with any $\psi : \mathbb{R}^d \rightarrow Z, \phi : Z \rightarrow Y$.
- ▶ $F(\mathbf{X}) = \text{rank}(\mathbf{X})$.

A formal definition of invariance

Function operating on sets/multisets

Let \mathcal{X} be a **countable set**. By construction, any function acting on sets $f : 2^{\mathcal{X}} \rightarrow Y$ for some Y is **permutation invariant**. That is

$$\forall \{x_1, \dots, x_n\} \in 2^{\mathcal{X}}, \forall \sigma \in S_n, f(\{x_1, \dots, x_n\}) = f(\{x_{\sigma^{-1}(1)}, \dots, x_{\sigma^{-1}(n)}\}).$$

Simply because $\{x_1, \dots, x_n\} = \{x_{\sigma^{-1}(1)}, \dots, x_{\sigma^{-1}(n)}\}$.

A formal definition of invariance

Function operating on sets/multisets

Let \mathcal{X} be a **countable set**. By construction, any function acting on sets $f : 2^{\mathcal{X}} \rightarrow Y$ for some Y is **permutation invariant**. That is

$$\forall \{x_1, \dots, x_n\} \in 2^{\mathcal{X}}, \forall \sigma \in S_n, f(\{x_1, \dots, x_n\}) = f(\{x_{\sigma^{-1}(1)}, \dots, x_{\sigma^{-1}(n)}\}).$$

Simply because $\{x_1, \dots, x_n\} = \{x_{\sigma^{-1}(1)}, \dots, x_{\sigma^{-1}(n)}\}$.

- ▶ Any function $f : 2^{\mathcal{X}} \rightarrow \mathbb{R}$ has the form ([Zaheer et al. 2018](#))

$$f(X) = \phi\left(\sum_{x \in X} \psi(x)\right) \text{ for some } \psi : \mathcal{X} \rightarrow \mathbb{R}, \phi : \mathbb{R} \rightarrow \mathbb{R}.$$

- ▶ See prev. course: a multiset is a “set” where element can be repeated several times e.g. $\{\{a, a, b\}\}$.
- ▶ Same representation result holds for functions on multisets ([Wagstaff et al. 2019](#)).

A formal definition of equivariance

Equivariance

Let Ω_1, Ω_2 be two \mathfrak{G} -sets (of the same group). A function $h : \Omega_1 \rightarrow \Omega_2$ is **\mathfrak{G} -equivariant** if

$$\forall x \in \Omega_1, \forall g \in \mathfrak{G}, h(gx) = gh(x).$$

- ▶ Pay attention to the input/output spaces and the compatibility.
- ▶ Transform the input + apply h = apply h and transform the result.

A formal definition of equivariance

Equivariance

Let Ω_1, Ω_2 be two \mathfrak{G} -sets (of the same group). A function $h : \Omega_1 \rightarrow \Omega_2$ is **\mathfrak{G} -equivariant** if

$$\forall x \in \Omega_1, \forall g \in \mathfrak{G}, h(gx) = gh(x).$$

- ▶ Pay attention to the input/output spaces and the compatibility.
- ▶ Transform the input + apply h = apply h and transform the result.

Convolutions

Prove that the convolution with a filter $h \in L_2(\mathbb{R})$ is translation equivariant.

A formal definition of equivariance

Equivariance

Let Ω_1, Ω_2 be two \mathfrak{G} -sets (of the same group). A function $h : \Omega_1 \rightarrow \Omega_2$ is **\mathfrak{G} -equivariant** if

$$\forall x \in \Omega_1, \forall g \in \mathfrak{G}, h(gx) = gh(x).$$

- ▶ Pay attention to the input/output spaces and the compatibility.
- ▶ Transform the input + apply h = apply h and transform the result.

Convolutions

Consider a filter $h \in L_2(\mathbb{R})$.

- ▶ The convolution with a filter is $H : \Omega = L_2(\mathbb{R}) \rightarrow L_2(\mathbb{R})$ such that $H(g) := g * h = h * g$.
- ▶ For any translation τ_x

$$\forall g \in L_2(\mathbb{R}), H(\tau_x g) = (\tau_x g) * h = \tau_x(g * h) = \tau_x H(g).$$

- ▶ Translate then convolve = convolve then translate.

A formal definition of equivariance

Equivariance

Let Ω_1, Ω_2 be two \mathfrak{G} -sets (of the same group). A function $h : \Omega_1 \rightarrow \Omega_2$ is **\mathfrak{G} -equivariant** if

$$\forall x \in \Omega_1, \forall g \in \mathfrak{G}, h(gx) = gh(x).$$

- ▶ Pay attention to the input/output spaces and the compatibility.
- ▶ Transform the input + apply h = apply h and transform the result.

Permutation equivariant functions

- ▶ Find two permutation equivariant functions $F : \mathbb{R}^{n \times d_1} \rightarrow \mathbb{R}^{n \times d_2}$.

A formal definition of equivariance

Equivariance

Let Ω_1, Ω_2 be two \mathfrak{G} -sets (of the same group). A function $h : \Omega_1 \rightarrow \Omega_2$ is **\mathfrak{G} -equivariant** if

$$\forall x \in \Omega_1, \forall g \in \mathfrak{G}, h(gx) = gh(x).$$

- ▶ Pay attention to the input/output spaces and the compatibility.
- ▶ Transform the input + apply h = apply h and transform the result.

Permutation equivariant functions

- ▶ Let $\mathbf{W} \in \mathbb{R}^{d_1 \times d_2}$ and $F(\mathbf{X}) = \mathbf{XW}$.

$$\text{▶ Let } \mathbf{X} = \begin{pmatrix} \mathbf{x}_1^\top \\ \vdots \\ \mathbf{x}_n^\top \end{pmatrix} \text{ previous example } F(\mathbf{X}) = \begin{pmatrix} (\mathbf{W}^\top \mathbf{x}_1)^\top \\ \vdots \\ (\mathbf{W}^\top \mathbf{x}_n)^\top \end{pmatrix}.$$

$$\text{▶ More generally } F(\mathbf{X}) = \begin{pmatrix} \psi(\mathbf{x}_1)^\top \\ \vdots \\ \psi(\mathbf{x}_n)^\top \end{pmatrix} \text{ where } \psi : \mathbb{R}^{d_1} \rightarrow \mathbb{R}^{d_2}.$$

A formal definition of equivariance

Equivariance

Let Ω_1, Ω_2 be two \mathfrak{G} -sets (of the same group). A function $h : \Omega_1 \rightarrow \Omega_2$ is **\mathfrak{G} -equivariant** if

$$\forall x \in \Omega_1, \forall g \in \mathfrak{G}, h(gx) = gh(x).$$

- ▶ Pay attention to the input/output spaces and the compatibility.
- ▶ Transform the input + apply h = apply h and transform the result.

Laplacian matrix

- ▶ An action of S_n on $\mathbb{R}^{n \times n}$ is defined as

$$\sigma \mathbf{A} = (A_{\sigma^{-1}(i), \sigma^{-1}(j)})_{ij}$$

- ▶ $\mathcal{L} : \text{sym}_n(\mathbb{R}) \rightarrow \text{sym}_n(\mathbb{R})$ which takes a symmetric matrix \mathbf{A} and outputs the Laplacian matrix $\mathcal{L}(\mathbf{A}) = \text{diag}(\mathbf{A}\mathbf{1}) - \mathbf{A}$
- ▶ **Show that \mathcal{L} is S_n -permutation equivariant.**

Combining them together

Composition of invariant/equivariant functions

Let Ω_1, Ω_2 be \mathfrak{G} -sets.

- ▶ Let $f : \Omega_1 \rightarrow \Omega_2$ be a **\mathfrak{G} -equivariant** function.
- ▶ Let $g : \Omega_2 \rightarrow Y$ be a **\mathfrak{G} -invariant** function.

Then $h = g \circ f$ is **\mathfrak{G} -invariant**.

Combining them together

Composition of invariant/equivariant functions

Let Ω_1, Ω_2 be \mathfrak{G} -sets.

- ▶ Let $f : \Omega_1 \rightarrow \Omega_2$ be a **\mathfrak{G} -equivariant** function.
- ▶ Let $g : \Omega_2 \rightarrow Y$ be a **\mathfrak{G} -invariant** function.

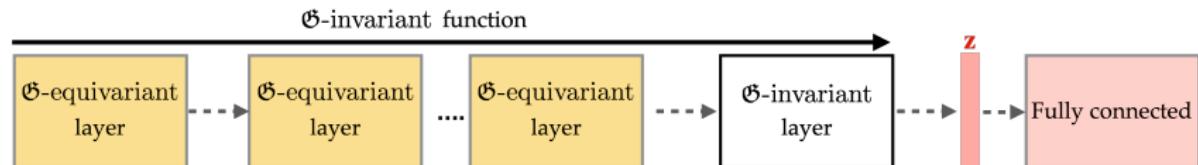
Then $h = g \circ f$ is **\mathfrak{G} -invariant**.

Proof

Indeed with $x \in \Omega_1, g \in \mathfrak{G}$

$$h(gx) = g(f(gx)) = g(gf(x)) = g(f(x)) = (g \circ f)(x) = h(x).$$

Simple but powerful: one of the reason CNNs work so well



Combining them together

Composition of invariant/equivariant functions

Let Ω_1, Ω_2 be \mathfrak{G} -sets.

- ▶ Let $f : \Omega_1 \times Y \rightarrow \Omega_2$ be a **\mathfrak{G} -equivariant** function **with respect to its first variable** i.e. $\forall y \in \Omega_1, \forall g \in \mathfrak{G}, \forall y \in Y, f(gx, y) = gf(x, y)$.
- ▶ Let $g : \Omega_1 \rightarrow Y$ be a **\mathfrak{G} -invariant** function.

Then the function h defined by $h(x) = f(x, g(x))$ is **\mathfrak{G} -equivariant**.

Proof

$$h(gx) = f(gx, g(gx)) = f(gx, g(x)) = gf(x, g(x)) = gh(x).$$

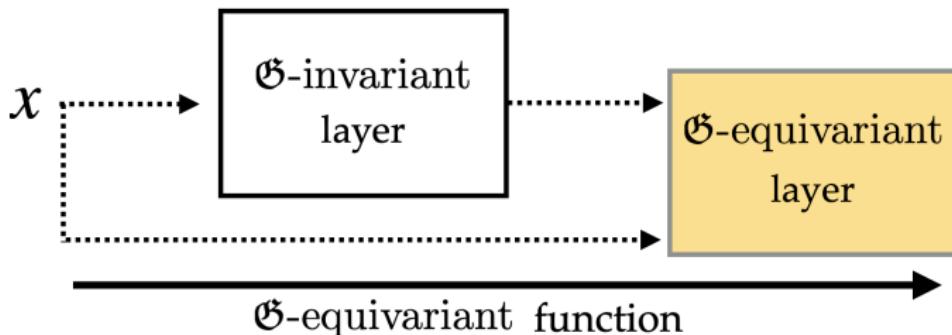


Table of contents

From neural networks...

The basic ideas

Logistic regression and one layer neural-network

Convolutional neural networks

Graph neural networks

Learning with graphs

What is a GNN ?

A bit of group theory

Invariance and equivariance

Permutation invariance/equivariance

Message-passing neural networks

Examples of GNN

The whole pipeline

Expressivity of GNN

Conclusion

More on unsupervised node embeddings techniques

Unsupervised node embeddings

Focus on permutation invariance/equivariance

Permutations as matrices

- ▶ $\sigma \in S_n$ can be described as $\mathbf{P}_\sigma = \begin{pmatrix} \mathbf{e}_{\sigma(1)}^\top \\ \vdots \\ \mathbf{e}_{\sigma(n)}^\top \end{pmatrix} \in \{0, 1\}^{n \times n}$. $\mathbf{P}_{\sigma^{-1}} = \mathbf{P}_\sigma^\top$.
- ▶ For $\mathbf{A} \in \mathbb{R}^{n \times n}$, the previous action is $\sigma\mathbf{A} = (A_{\sigma^{-1}(i)\sigma^{-1}(j)})_{ij} = \mathbf{P}_\sigma^\top \mathbf{A} \mathbf{P}_\sigma$.
- ▶ An action of S_n on $\mathbb{R}^{n \times d} \times \mathbb{R}^{n \times n}$

$$\sigma \cdot (\mathbf{X}, \mathbf{A}) = (\mathbf{P}_\sigma^\top \mathbf{X}, \mathbf{P}_\sigma^\top \mathbf{A} \mathbf{P}_\sigma)$$

Focus on permutation invariance/equivariance

Permutations as matrices

- ▶ $\sigma \in S_n$ can be described as $\mathbf{P}_\sigma = \begin{pmatrix} \mathbf{e}_{\sigma(1)}^\top \\ \vdots \\ \mathbf{e}_{\sigma(n)}^\top \end{pmatrix} \in \{0, 1\}^{n \times n}$. $\mathbf{P}_{\sigma^{-1}} = \mathbf{P}_\sigma^\top$.
- ▶ For $\mathbf{A} \in \mathbb{R}^{n \times n}$, the previous action is $\sigma \mathbf{A} = (A_{\sigma^{-1}(i)\sigma^{-1}(j)})_{ij} = \mathbf{P}_\sigma^\top \mathbf{A} \mathbf{P}_\sigma$.
- ▶ An action of S_n on $\mathbb{R}^{n \times d} \times \mathbb{R}^{n \times n}$

$$\sigma \cdot (\mathbf{X}, \mathbf{A}) = (\mathbf{P}_\sigma^\top \mathbf{X}, \mathbf{P}_\sigma^\top \mathbf{A} \mathbf{P}_\sigma)$$

Interpretation

- ▶ $\sigma \cdot (\mathbf{X}, \mathbf{A})$ permutes the nodes of the graph **and** the features **in the same manner**.

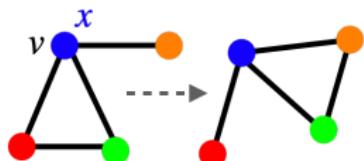


Figure: Is it a valid action of σ ?

Focus on permutation invariance/equivariance

Permutations as matrices

- ▶ $\sigma \in S_n$ can be described as $\mathbf{P}_\sigma = \begin{pmatrix} \mathbf{e}_{\sigma(1)}^\top \\ \vdots \\ \mathbf{e}_{\sigma(n)}^\top \end{pmatrix} \in \{0, 1\}^{n \times n}$. $\mathbf{P}_{\sigma^{-1}} = \mathbf{P}_\sigma^\top$.
- ▶ For $\mathbf{A} \in \mathbb{R}^{n \times n}$, the previous action is $\sigma \mathbf{A} = (A_{\sigma^{-1}(i)\sigma^{-1}(j)})_{ij} = \mathbf{P}_\sigma^\top \mathbf{A} \mathbf{P}_\sigma$.
- ▶ An action of S_n on $\mathbb{R}^{n \times d} \times \mathbb{R}^{n \times n}$

$$\sigma \cdot (\mathbf{X}, \mathbf{A}) = (\mathbf{P}_\sigma^\top \mathbf{X}, \mathbf{P}_\sigma^\top \mathbf{A} \mathbf{P}_\sigma)$$

Back to the GNN context

- ▶ In classification/regression $f : G = (\mathbf{X}, \mathbf{A}) \rightarrow y \in Y$ (e.g. $(\{+1, -1\})$).
- ▶ For node embeddings $F : G = (\mathbf{X}, \mathbf{A}) \rightarrow \mathbf{Z} \in \mathbb{R}^{n \times k}$

Focus on permutation invariance/equivariance

Permutations as matrices

- ▶ $\sigma \in S_n$ can be described as $\mathbf{P}_\sigma = \begin{pmatrix} \mathbf{e}_{\sigma(1)}^\top \\ \vdots \\ \mathbf{e}_{\sigma(n)}^\top \end{pmatrix} \in \{0, 1\}^{n \times n}$. $\mathbf{P}_{\sigma^{-1}} = \mathbf{P}_\sigma^\top$.
- ▶ For $\mathbf{A} \in \mathbb{R}^{n \times n}$, the previous action is $\sigma \mathbf{A} = (A_{\sigma^{-1}(i)\sigma^{-1}(j)})_{ij} = \mathbf{P}_\sigma^\top \mathbf{A} \mathbf{P}_\sigma$.
- ▶ An action of S_n on $\mathbb{R}^{n \times d} \times \mathbb{R}^{n \times n}$

$$\sigma \cdot (\mathbf{X}, \mathbf{A}) = (\mathbf{P}_\sigma^\top \mathbf{X}, \mathbf{P}_\sigma^\top \mathbf{A} \mathbf{P}_\sigma)$$

Back to the GNN context

- ▶ In classification/regression $f : G = (\mathbf{X}, \mathbf{A}) \rightarrow y \in Y$ (e.g. $(\{+1, -1\})$).
- ▶ For node embeddings $F : G = (\mathbf{X}, \mathbf{A}) \rightarrow \mathbf{Z} \in \mathbb{R}^{n \times k}$

Ensuring invariance/equivariance is key when learning on graphs

Find f that are S_n -**invariant**, F that are S_n -**equivariant**.

- ▶ $f(\mathbf{P}_\sigma^\top \mathbf{X}, \mathbf{P}_\sigma^\top \mathbf{A} \mathbf{P}_\sigma) = f(\mathbf{X}, \mathbf{A})$ and $F(\mathbf{P}_\sigma^\top \mathbf{X}, \mathbf{P}_\sigma^\top \mathbf{A} \mathbf{P}_\sigma) = \mathbf{P}_\sigma^\top F(\mathbf{X}, \mathbf{A})$.

Focus on permutation invariance/equivariance

Ensuring invariance/equivariance is key when learning on graphs

Find f that are S_n -**invariant**, F that are S_n -**equivariant**.

- ▶ $f(\mathbf{P}_\sigma^\top \mathbf{X}, \mathbf{P}_\sigma^\top \mathbf{A} \mathbf{P}_\sigma) = f(\mathbf{X}, \mathbf{A})$ and $F(\mathbf{P}_\sigma^\top \mathbf{X}, \mathbf{P}_\sigma^\top \mathbf{A} \mathbf{P}_\sigma) = \mathbf{P}_\sigma^\top F(\mathbf{X}, \mathbf{A})$.

Examples: equivariance (1/2)

- ▶ Take $\mathbf{X} \in \mathbb{R}^{n \times d_1}$, $\mathbf{W} \in \mathbb{R}^{d_1 \times d_2}$ and a function Ψ that applies **independently on each row** of a matrix.
- ▶ $F(\mathbf{X}, \mathbf{A}) = \Psi(\mathbf{A} \mathbf{X} \mathbf{W})$ is S_n -**equivariant**.

Focus on permutation invariance/equivariance

Ensuring invariance/equivariance is key when learning on graphs

Find f that are S_n -**invariant**, F that are S_n -**equivariant**.

- ▶ $f(\mathbf{P}_\sigma^\top \mathbf{X}, \mathbf{P}_\sigma^\top \mathbf{A} \mathbf{P}_\sigma) = f(\mathbf{X}, \mathbf{A})$ and $F(\mathbf{P}_\sigma^\top \mathbf{X}, \mathbf{P}_\sigma^\top \mathbf{A} \mathbf{P}_\sigma) = \mathbf{P}_\sigma^\top F(\mathbf{X}, \mathbf{A})$.

Examples: equivariance (1/2)

- ▶ Take $\mathbf{X} \in \mathbb{R}^{n \times d_1}$, $\mathbf{W} \in \mathbb{R}^{d_1 \times d_2}$ and a function Ψ that applies **independently on each row** of a matrix.
- ▶ $F(\mathbf{X}, \mathbf{A}) = \Psi(\mathbf{A}\mathbf{X}\mathbf{W})$ is S_n -**equivariant**.
- ▶ In particular when Ψ is **element-wise**.

Focus on permutation invariance/equivariance

Ensuring invariance/equivariance is key when learning on graphs

Find f that are S_n -**invariant**, F that are S_n -**equivariant**.

- ▶ $f(\mathbf{P}_\sigma^\top \mathbf{X}, \mathbf{P}_\sigma^\top \mathbf{A} \mathbf{P}_\sigma) = f(\mathbf{X}, \mathbf{A})$ and $F(\mathbf{P}_\sigma^\top \mathbf{X}, \mathbf{P}_\sigma^\top \mathbf{A} \mathbf{P}_\sigma) = \mathbf{P}_\sigma^\top F(\mathbf{X}, \mathbf{A})$.

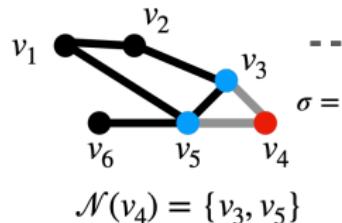
Examples: equivariance (1/2)

- ▶ Take $\mathbf{X} \in \mathbb{R}^{n \times d_1}$, $\mathbf{W} \in \mathbb{R}^{d_1 \times d_2}$ and a function Ψ that applies **independently on each row** of a matrix.
- ▶ $F(\mathbf{X}, \mathbf{A}) = \Psi(\mathbf{A}\mathbf{X}\mathbf{W})$ is S_n -**equivariant**.
- ▶ In particular when Ψ is **element-wise**.
- ▶ But also $F(\mathbf{X}, \mathbf{A}) = \Psi(G(\mathbf{A})\mathbf{X}\mathbf{W})$ where G is S_n -**equivariant**.
- ▶ E.g. $F(\mathbf{X}, \mathbf{A}) = \Psi(\mathcal{L}(\mathbf{A})\mathbf{X}\mathbf{W})$ where \mathcal{L} computes the Laplacian.
- ▶ E.g. $F(\mathbf{X}, \mathbf{A}) = \Psi(P[\mathcal{L}](\mathbf{A})\mathbf{X}\mathbf{W})$ where P is a polynomial $P[\mathcal{L}] = \sum_m c_m \mathcal{L}^m$.

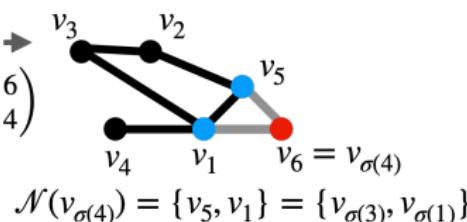
Focus on permutation invariance/equivariance

Examples: equivariance (2/2)

- Take $\mathbf{X} = \begin{pmatrix} \mathbf{x}_1^\top \\ \vdots \\ \mathbf{x}_n^\top \end{pmatrix}$ and define the multiset $X_i := \{\{\mathbf{x}_j : j \in \mathcal{N}(i)\}\}$.
- Then $X_{\sigma(i)} = \{\{\mathbf{x}_{\sigma(j)} : j \in \mathcal{N}(i)\}\}$:



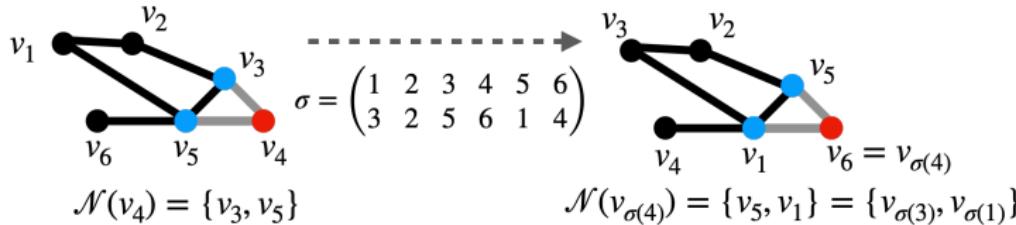
$$\sigma = \begin{pmatrix} 1 & 2 & 3 & 4 & 5 & 6 \\ 3 & 2 & 5 & 6 & 1 & 4 \end{pmatrix}$$



Focus on permutation invariance/equivariance

Examples: equivariance (2/2)

- Take $\mathbf{X} = \begin{pmatrix} \mathbf{x}_1^\top \\ \vdots \\ \mathbf{x}_n^\top \end{pmatrix}$ and define the multiset $X_i := \{\{\mathbf{x}_j : j \in \mathcal{N}(i)\}\}$.
- Then $X_{\sigma(i)} = \{\{\mathbf{x}_{\sigma(j)} : j \in \mathcal{N}(i)\}\}$:



- A function AGGREGATE operating on **multisets of vectors**.
- Then the following function is permutation equivariant.

$$F(\mathbf{X}, \mathbf{A}) = \begin{pmatrix} \psi(\mathbf{x}_1, \text{AGGREGATE}(X_1)) \\ \vdots \\ \psi(\mathbf{x}_n, \text{AGGREGATE}(X_n)) \end{pmatrix}$$

Table of contents

From neural networks...

The basic ideas

Logistic regression and one layer neural-network

Convolutional neural networks

Graph neural networks

Learning with graphs

What is a GNN ?

A bit of group theory

Invariance and equivariance

Permutation invariance/equivariance

Message-passing neural networks

Examples of GNN

The whole pipeline

Expressivity of GNN

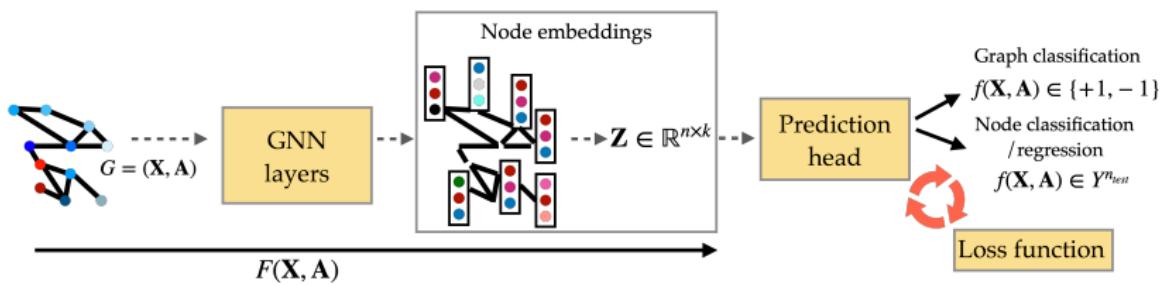
Conclusion

More on unsupervised node embeddings techniques

Unsupervised node embeddings

Remember

The training pipeline

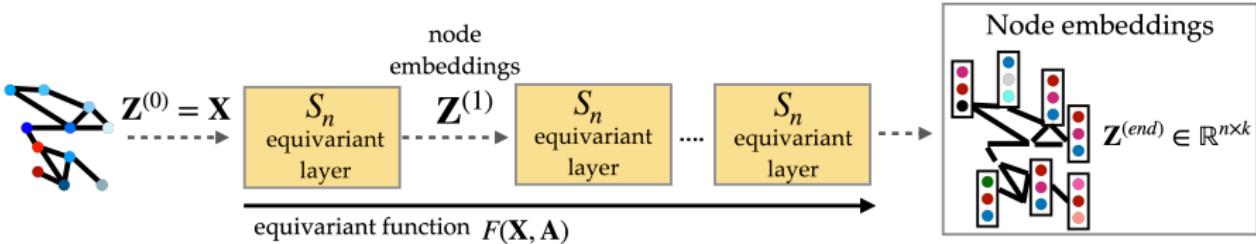


- ▶ Find an embedding of the nodes $F(\mathbf{X}, \mathbf{A}) \in \mathbb{R}^{n \times k}$ and then do stuff.

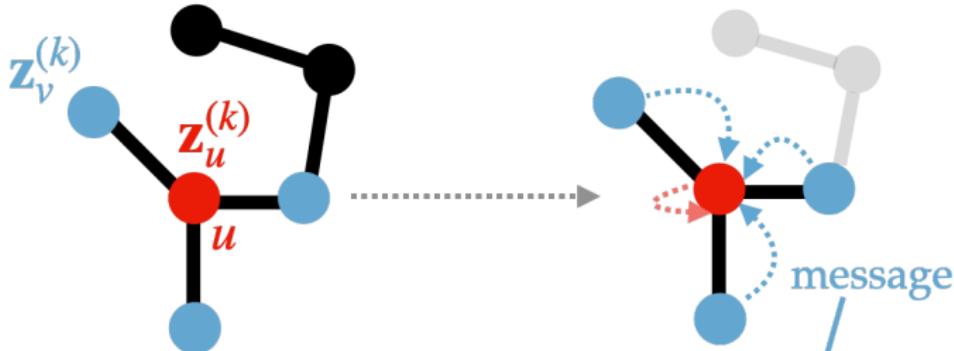
Message-passing for node embeddings

Goal of the message passing framework

- ▶ Defines specific S_n -equivariant layers/functions.
- ▶ Can be used for node embeddings.
- ▶ Usually $\mathbf{Z}^{(0)} = \mathbf{X}$ but when no node features are available several options (e.g. node statistics).
- ▶ Notation: $\mathbf{z}_u^{(k)}$ is the embedding of the node $u \in V$ at the k -layer.



The message passing framework



$$\mathbf{z}_u^{(k+1)} = \text{COMBINE}^{(k)} \left(\mathbf{z}_u^{(k)}, \text{AGGREGATE}^{(k)} \left(\{\{\mathbf{z}_v^{(k)} : v \in \mathcal{N}(u)\}\} \right) \right)$$

One of the most used GNN framework in practice

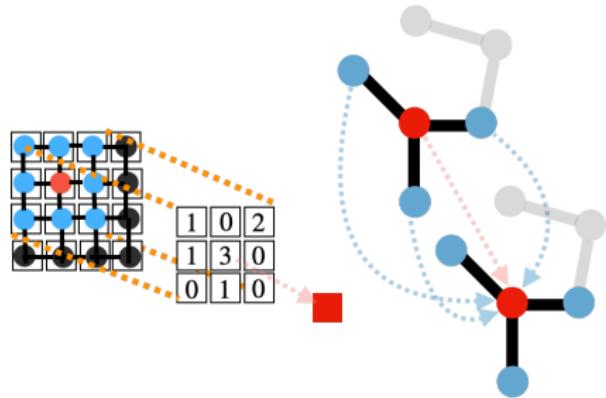
- ▶ At each iteration, **every node aggregates information from its local neighborhood.**
- ▶ A zoo of methods for different COMBINE, AGGREGATE functions.
- ▶ Why is this defining a permutation equivariant layer ?

The message passing framework

Similarities with CNN

- ▶ One layer of message-passing GNN shares similarities to convolutional layers.
- ▶ Usually it takes the form

$$\mathbf{z}_u^{(k+1)} = \phi \left(\sum_{v \in \mathcal{N}(u) \cup \{u\}} \alpha_{uv} \mathbf{z}_v^{(k)} \right).$$

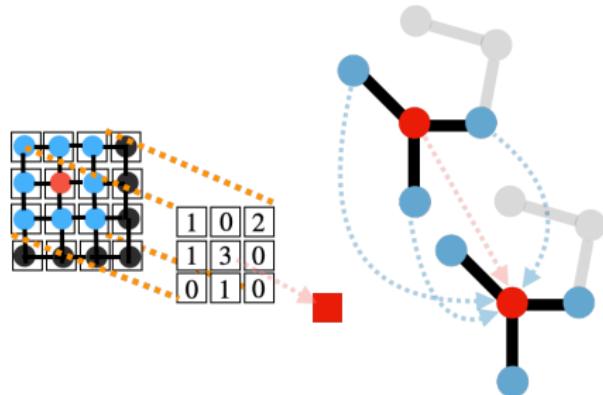


The message passing framework

Similarities with CNN

- ▶ One layer of message-passing GNN shares similarities to convolutional layers.
- ▶ Usually it takes the form

$$\mathbf{z}_u^{(k+1)} = \phi \left(\sum_{v \in \mathcal{N}(u) \cup \{u\}} \alpha_{uv} \mathbf{z}_v^{(k)} \right).$$



k-hop neighbourhood

After k -steps each node has received the informations from its k -hop neighbourhood.

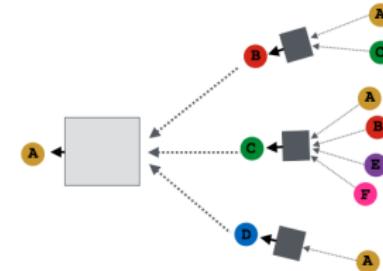
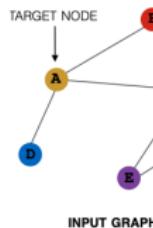


Figure: From Jure Leskovec course *Machine Learning with Graphs*.

Table of contents

From neural networks...

The basic ideas

Logistic regression and one layer neural-network

Convolutional neural networks

Graph neural networks

Learning with graphs

What is a GNN ?

A bit of group theory

Invariance and equivariance

Permutation invariance/equivariance

Message-passing neural networks

Examples of GNN

The whole pipeline

Expressivity of GNN

Conclusion

More on unsupervised node embeddings techniques

Unsupervised node embeddings

A first GNN with message passing

Sum/mean aggregation (Scarselli et al. 2008)

A first idea would be

$$\mathbf{z}_u^{(k+1)} = \phi(\mathbf{W}_{\text{self}}^{(k)} \mathbf{z}_u^{(k)} + \mathbf{W}_{\text{neigh}}^{(k)} \sum_{v \in \mathcal{N}(u)} \mathbf{z}_v^{(k)} + \mathbf{b}^{(k)})$$

- ▶ $\mathbf{W}_{\text{self}}^{(k)}, \mathbf{W}_{\text{neigh}}^{(k)} \in \mathbb{R}^{d_{k+1} \times d_k}$ are matrices of **learnable parameters**.
- ▶ Do not depend on the number of nodes ! .
- ▶ Complexity of computing it for all nodes is $O(|E|)$.
- ▶ $\mathbf{b}^{(k)} \in \mathbb{R}^{d_{k+1}}$ is a bias term (often omitted to simplify notations).
- ▶ ϕ is a **pointwise** non-linearity such as ReLu.

Questions

- ▶ What is COMBINE, AGGREGATE ?
- ▶ Write this in matrix form.

A first GNN with message passing

Sum/mean aggregation (Scarselli et al. 2008)

A first idea would be

$$\mathbf{z}_u^{(k+1)} = \phi(\mathbf{W}_{\text{self}}^{(k)} \mathbf{z}_u^{(k)} + \mathbf{W}_{\text{neigh}}^{(k)} \sum_{v \in \mathcal{N}(u)} \mathbf{z}_v^{(k)} + \mathbf{b}^{(k)})$$

- ▶ $\mathbf{W}_{\text{self}}^{(k)}, \mathbf{W}_{\text{neigh}}^{(k)} \in \mathbb{R}^{d_{k+1} \times d_k}$ are matrices of **learnable parameters**.
- ▶ Do not depend on the number of nodes ! .
- ▶ Complexity of computing it for all nodes is $O(|E|)$.
- ▶ $\mathbf{b}^{(k)} \in \mathbb{R}^{d_{k+1}}$ is a bias term (often omitted to simplify notations).
- ▶ ϕ is a **pointwise** non-linearity such as ReLu.

Answers

- ▶ What is COMBINE, AGGREGATE ?
- ▶ $\forall k, \text{AGGREGATE}^{(k)}(\{\{\mathbf{z}_v : v \in \mathcal{N}(u)\}\}) = \sum_{v \in \mathcal{N}(u)} \mathbf{z}_v$.
- ▶ $\text{COMBINE}^{(k)}(\mathbf{z}_1, \mathbf{z}_2) = \phi(\mathbf{W}_{\text{self}}^{(k)} \mathbf{z}_1 + \mathbf{W}_{\text{neigh}}^{(k)} \mathbf{z}_2 + \mathbf{b}^{(k)})$.

A first GNN with message passing

Sum/mean aggregation (Scarselli et al. 2008)

A first idea would be

$$\mathbf{z}_u^{(k+1)} = \phi(\mathbf{W}_{\text{self}}^{(k)} \mathbf{z}_u^{(k)} + \mathbf{W}_{\text{neigh}}^{(k)} \sum_{v \in \mathcal{N}(u)} \mathbf{z}_v^{(k)} + \mathbf{b}^{(k)})$$

- ▶ $\mathbf{W}_{\text{self}}^{(k)}, \mathbf{W}_{\text{neigh}}^{(k)} \in \mathbb{R}^{d_{k+1} \times d_k}$ are matrices of **learnable parameters**.
- ▶ Do not depend on the number of nodes ! .
- ▶ Complexity of computing it for all nodes is $O(|E|)$.
- ▶ $\mathbf{b}^{(k)} \in \mathbb{R}^{d_{k+1}}$ is a bias term (often omitted to simplify notations).
- ▶ ϕ is a **pointwise** non-linearity such as ReLu.

Answers

- ▶ Write this in matrix form.

$$\mathbf{Z}^{(k+1)} = \phi \left(\mathbf{A} \mathbf{Z}^{(k)} \mathbf{W}_{\text{neigh}}^{(k)} + \mathbf{Z}^{(k)} \mathbf{W}_{\text{self}}^{(k)} + \begin{pmatrix} \mathbf{b}^{(k)} \\ \vdots \\ \mathbf{b}^{(k)} \end{pmatrix} \right).$$

Graph convolutional neural networks

Most popular baseline model

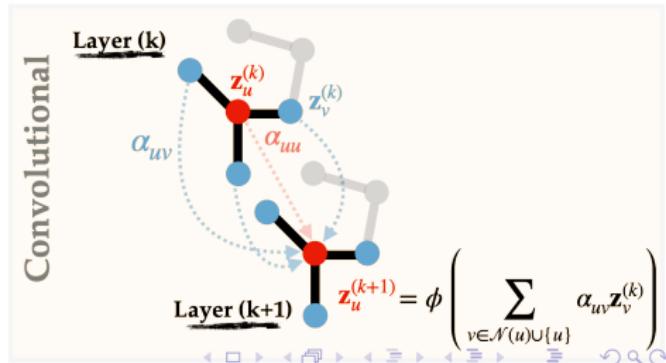
Introduced by Kipf and Welling 2016 for semi-supervised node classification.

$$\mathbf{z}_u^{(k+1)} = \text{Relu}(\mathbf{W}_{\text{self}}^{(k)} \mathbf{z}_u^{(k)} + \mathbf{W}_{\text{neigh}}^{(k)} \frac{1}{\sqrt{|\mathcal{N}(u)|}} \sum_{v \in \mathcal{N}(u)} \frac{\mathbf{z}_v^{(k)}}{\sqrt{|\mathcal{N}(v)|}})$$

- ▶ Also GraphSage framework (William L. Hamilton, R. Ying, and Leskovec 2018).
- ▶ What is COMBINE, AGGREGATE ?

In matrix form

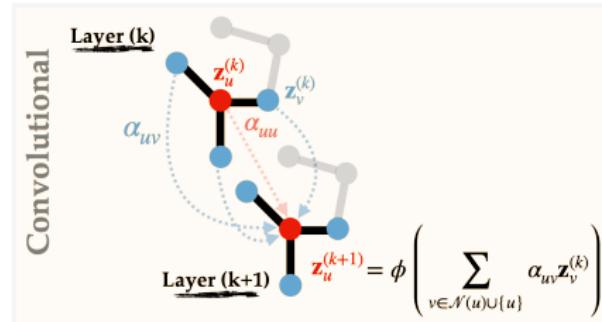
- ▶ With $\mathbf{W}_{\text{self}} = \mathbf{W}_{\text{neigh}}$, $\mathbf{Z}^{(k+1)} = \text{Relu} \left((\mathbf{I} + \mathbf{D}^{-\frac{1}{2}} \mathbf{A} \mathbf{D}^{-\frac{1}{2}}) \mathbf{Z}^{(k)} \mathbf{W}^{(k)} \right)$.
- ▶ First-order approximation of localized spectral filters on graphs.



Graph Attention Networks

Motivations

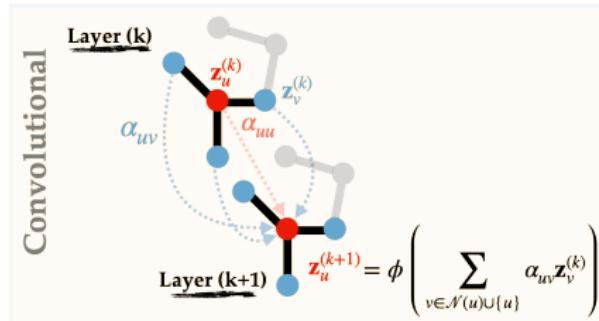
- ▶ In many MP-GNN layers weights of the convolutions **are fixed**.
- ▶ What if we also learn them ?
- ▶ Learn the importance of the neighbours contributions.



Graph Attention Networks

Motivations

- ▶ In many MP-GNN layers weights of the convolutions **are fixed**.
- ▶ What if we also learn them ?
- ▶ Learn the importance of the neighbours contributions.



GAT networks (Veličković et al. 2017)

$$z_u^{(k+1)} = \text{Relu}(\mathbf{W}^{(k)} \sum_{v \in \mathcal{N}(u) \cup \{u\}} \alpha_{uv} z_v^{(k)})$$

- ▶ Here α_{uv} are **learnable weights**.
- ▶ $e_{uv} = \text{NN}(\Theta_1 z_u, \Theta_2 z_u)$ with learnable matrices Θ_1, Θ_2 and

$$\alpha_{uv} = \text{softmax}_v(e_{uv}) = \frac{\exp(e_{uv})}{\sum_{v' \in \mathcal{N}(u)} e_{uv'}}$$

- ▶ It is based on attention mechanisms (Vaswani et al. 2023).

Graph Isomorphism Networks (GIN)

The problem of injectivity

Xu et al. 2019 provide a detailed discussion of the relative power of GNN.

- ▶ One interesting property is **injectivity** of COMBINE, AGGREGATE.
- ▶ They propose

$$\mathbf{z}_u^{(k+1)} = \text{MLP}^{(k)} \left((1 + \theta^{(k)}) \mathbf{z}_u^{(k)} + \sum_{v \in \mathcal{N}(u)} \mathbf{z}_v^{(k)} \right)$$

- ▶ $\text{MLP} : \mathbb{R}^{d_k} \rightarrow \mathbb{R}^{d_{k+1}}$ is a fully connected neural-network.

Spectral GNN

Learning filters

Originally introduced by Bruna et al. 2013. The idea is

$$\mathbf{Z}^{(k+1)} = \text{Relu}(P[\mathcal{L}](\mathbf{A})\mathbf{Z}^{(k)}\mathbf{W}^{(k)})$$

- ▶ $\mathcal{L}(\mathbf{A}) = \text{diag}(\mathbf{A}\mathbf{1}) - \mathbf{A}$ is the Laplacian (or normalized version).
- ▶ $P[\mathcal{L}] = \sum_{m=0}^M c_m \mathcal{L}^m$ is a **learnable** polynomial of the Laplacian.
- ▶ As $\mathcal{L}(\mathbf{A}) = \mathbf{U}\Lambda\mathbf{U}^\top$, $P[\mathcal{L}](\mathbf{A}) = \mathbf{U}P[\Lambda]\mathbf{U}^\top$.
- ▶ Connections with the Fourier transform on graphs: $P[\mathcal{L}]$ acts as a filter.

Spectral GNN

Learning filters

Originally introduced by Bruna et al. 2013. The idea is

$$\mathbf{Z}^{(k+1)} = \text{Relu}(P[\mathcal{L}](\mathbf{A})\mathbf{Z}^{(k)}\mathbf{W}^{(k)})$$

- ▶ $\mathcal{L}(\mathbf{A}) = \text{diag}(\mathbf{A}\mathbf{1}) - \mathbf{A}$ is the Laplacian (or normalized version).
- ▶ $P[\mathcal{L}] = \sum_{m=0}^M c_m \mathcal{L}^m$ is a **learnable** polynomial of the Laplacian.
- ▶ As $\mathcal{L}(\mathbf{A}) = \mathbf{U}\Lambda\mathbf{U}^\top$, $P[\mathcal{L}](\mathbf{A}) = \mathbf{U}P[\Lambda]\mathbf{U}^\top$.
- ▶ Connections with the Fourier transform on graphs: $P[\mathcal{L}]$ acts as a filter.

Limitations

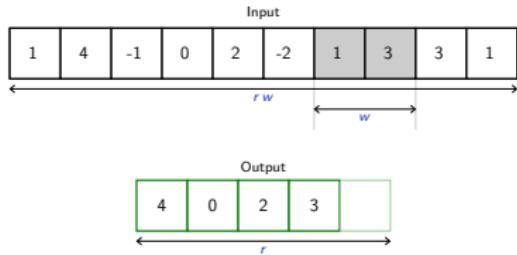
- ▶ Naive complexity in $O(|V|^3)$ (eigen-decomposition).
- ▶ Any perturbation to a graph results in a change of eigenbasis \mathbf{U} .
- ▶ Learned filters are domain dependent.
- ▶ Alternative ChebNet [Defferrard, Bresson, and Vandergheynst 2017](#) relies on Chebyshev polynomials with $O(|E|M)$ complexity.

Graph pooling

Pooling layers in neural networks

At the core of many NN architectures.

- ▶ Most standard type is **max-pooling**.
- ▶ ↓ the number of parameters to learn.
- ▶ Improves robustness.

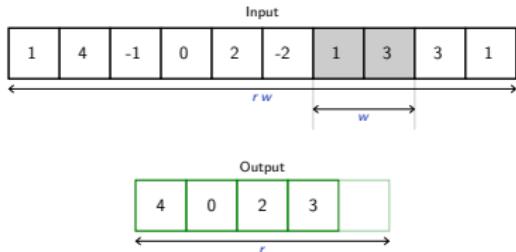


Graph pooling

Pooling layers in neural networks

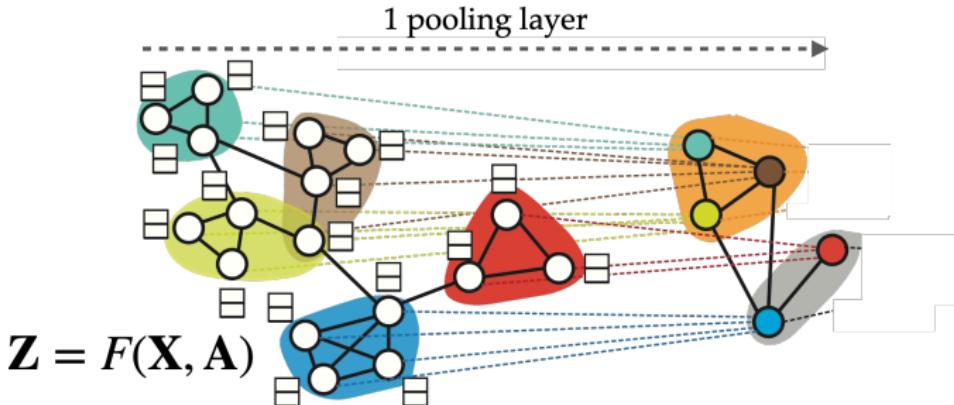
At the core of many NN architectures.

- ▶ Most standard type is **max-pooling**.
- ▶ ↓ the number of parameters to learn.
- ▶ Improves robustness.



Pooling in GNN

Equivalent to down-sampling = reducing the number of nodes.



Learning at the graph level

- ▶ The neural message passing approach produces a set of node embeddings $F(\mathbf{X}, \mathbf{A}) = \mathbf{Z} \in \mathbb{R}^{n \times k}$.
- ▶ What about predictions at the graph level ? E.g. in graph classification.
- ▶ **We want one embedding for the entire graph \mathbf{z}_G .**
- ▶ It should be a permutation invariant function $f(\mathbf{X}, \mathbf{A})$.
- ▶ E.g. global average pooling $\mathbf{z}_G = f(\mathbf{X}, \mathbf{A}) = \frac{1}{|V|} \sum_{u \in V} \mathbf{z}_u \in \mathbb{R}^k$.

Diffpool

Learning at the graph level

- ▶ The neural message passing approach produces a set of node embeddings $F(\mathbf{X}, \mathbf{A}) = \mathbf{Z} \in \mathbb{R}^{n \times k}$.
- ▶ What about predictions at the graph level ? E.g. in graph classification.
- ▶ **We want one embedding for the entire graph \mathbf{z}_G .**
- ▶ It should be a permutation invariant function $f(\mathbf{X}, \mathbf{A})$.
- ▶ E.g. global average pooling $\mathbf{z}_G = f(\mathbf{X}, \mathbf{A}) = \frac{1}{|V|} \sum_{u \in V} \mathbf{z}_u \in \mathbb{R}^k$.

Better idea: hierarchical pooling (Z. Ying et al. 2018)

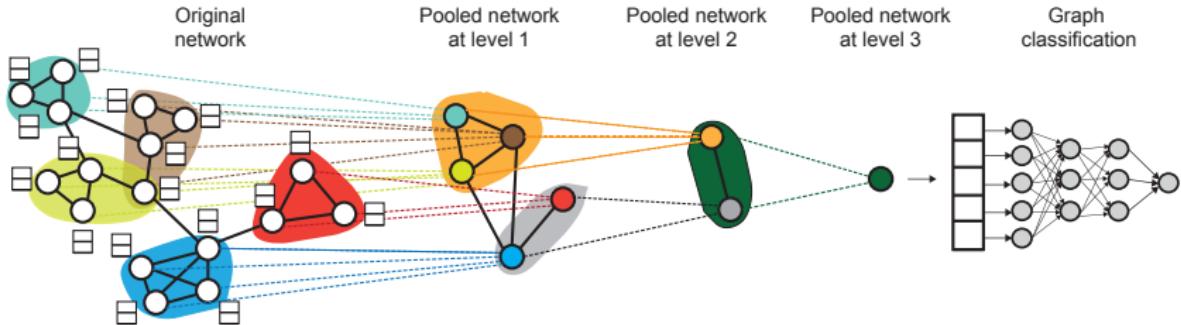


Table of contents

From neural networks...

The basic ideas

Logistic regression and one layer neural-network

Convolutional neural networks

Graph neural networks

Learning with graphs

What is a GNN ?

A bit of group theory

Invariance and equivariance

Permutation invariance/equivariance

Message-passing neural networks

Examples of GNN

The whole pipeline

Expressivity of GNN

Conclusion

More on unsupervised node embeddings techniques

Unsupervised node embeddings

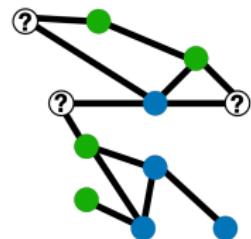
Applications

Node classification

- ▶ **One graph** G where each node has a class.

Train GNNs in a fully-supervised manner by minimizing

$$\mathcal{L} = \text{CE} = \sum_{u \in V_{train}} \sum_{k \in [K]} -[\mathbf{y}_u]_k \log([\text{softmax}(\mathbf{z}_u)]_k)$$



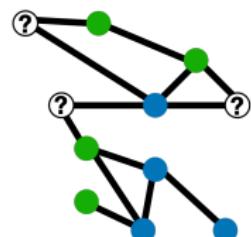
Applications

Node classification

- **One graph** G where each node has a class.

Train GNNs in a fully-supervised manner by minimizing

$$\mathcal{L} = \text{CE} = \sum_{u \in V_{\text{train}}} \sum_{k \in [K]} -[\mathbf{y}_u]_k \log([\text{softmax}(\mathbf{z}_u)]_k)$$



Graph classification

- **Many graphs** G_1, \dots, G_n associated with classes $(y_{G_i})_i$.

Train GNNs in a fully-supervised manner by minimizing

$$\mathcal{L} = \sum_{G \in T_{\text{train}}} \ell(\text{MLP}(\mathbf{z}_G), y_G)$$

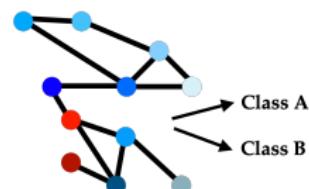


Table of contents

From neural networks...

The basic ideas

Logistic regression and one layer neural-network

Convolutional neural networks

Graph neural networks

Learning with graphs

What is a GNN ?

A bit of group theory

Invariance and equivariance

Permutation invariance/equivariance

Message-passing neural networks

Examples of GNN

The whole pipeline

Expressivity of GNN

Conclusion

More on unsupervised node embeddings techniques

Unsupervised node embeddings

Connection with the WL test

WL algorithm and MP-GNN

- ▶ WL algorithm and the message passing GNN approach are very similar.
- ▶ Iteratively aggregate information from local node neighborhoods.

Connection with the WL test

WL algorithm and MP-GNN

- ▶ WL algorithm and the message passing GNN approach are very similar.
- ▶ Iteratively aggregate information from local node neighborhoods.

Message passing neural networks are not that powerful ?

- ▶ Consider a MP-GNN with K layers

$$\mathbf{z}_u^{(k+1)} = \text{COMBINE}^{(k)} \left(\mathbf{z}_u^{(k)}, \text{AGGREGATE}^{(k)} \left(\{\{\mathbf{z}_v^{(k)} : v \in \mathcal{N}(u)\}\} \right) \right)$$

- ▶ Suppose that discrete node labels $\mathbf{Z}^{(0)} = \mathbf{X} \in \mathbb{Z}^{n \times d}$.
- ▶ Then Xu et al. 2019 show that

$$\mathbf{z}_u^{(K)} \neq \mathbf{z}_v^{(K)} \iff \text{labels of } u \text{ and } v \text{ are } \neq \text{ after } K \text{ iter. of the WL algorithm.}$$

- ▶ If the WL test cannot distinguish between G_1, G_2 , then MP-GNN also incapable of doing it.
- ▶ Ability of solving isomorphism = good measure of “expressivity” ?

Other limitations

- ▶ **The oversmoothing problem:** if too many layers of MP-GNN, the node features tend to converge to a non-informative limit.

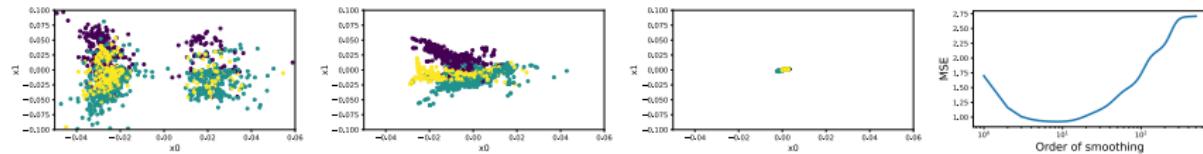


Figure: From Keriven 2022

- ▶ **Heterophily vs homophilie:** neighbours should have similar embeddings ? ([Luan et al. 2022](#)).

Table of contents

From neural networks...

The basic ideas

Logistic regression and one layer neural-network

Convolutional neural networks

Graph neural networks

Learning with graphs

What is a GNN ?

A bit of group theory

Invariance and equivariance

Permutation invariance/equivariance

Message-passing neural networks

Examples of GNN

The whole pipeline

Expressivity of GNN

Conclusion

More on unsupervised node embeddings techniques

Unsupervised node embeddings

Conclusion

- ▶ Flexible: graph/node/edge classification, semi-supervised learning, link prediction...
- ▶ Generally state-of-the-art, but...
- ▶ ... sometimes do not work “that well” (compared to other DL)
- ▶ Simple methods may perform better but might be “forgotten” in benchmarks
- ▶ Room for improvement (many interesting challenges), but conventional DL wisdom might not hold
- ▶ Arguably, no real “ImageNet moment” yet for GNNs → several recent initiatives for bigger datasets and more complex tasks (eg Open Graph Benchmark)

Table of contents

From neural networks...

The basic ideas

Logistic regression and one layer neural-network

Convolutional neural networks

Graph neural networks

Learning with graphs

What is a GNN ?

A bit of group theory

Invariance and equivariance

Permutation invariance/equivariance

Message-passing neural networks

Examples of GNN

The whole pipeline

Expressivity of GNN

Conclusion

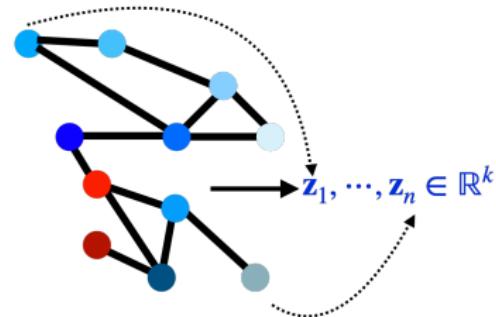
More on unsupervised node embeddings techniques

Unsupervised node embeddings

Objective

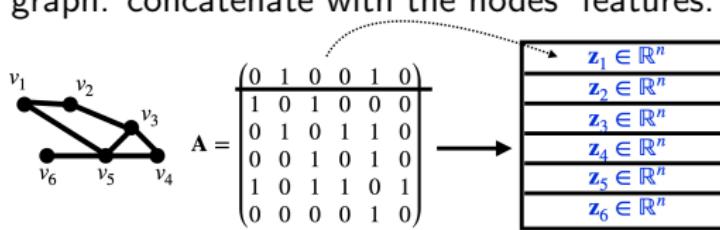
A chronological start

- ▶ Idea: to learn on a graph: nodes → vector → standard ML pipeline.
- ▶ The embedding **must take into account the structure** of the graph.
- ▶ Also useful for visualization.



One naive approach

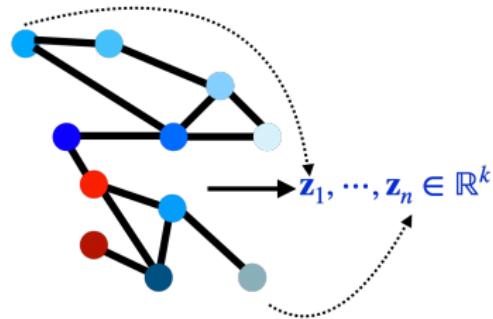
- ▶ Consider each row of the adjacency matrix as an embedding vector.
- ▶ If labelled graph: concatenate with the nodes' features.



Objective

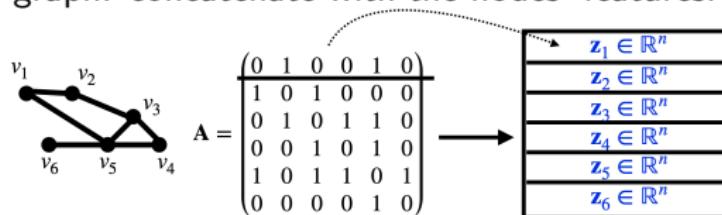
A chronological start

- ▶ Idea: to learn on a graph: nodes → vector → standard ML pipeline.
- ▶ The embedding **must take into account the structure** of the graph.
- ▶ Also useful for visualization.



One naive approach

- ▶ Consider each row of the adjacency matrix as an embedding vector.
- ▶ If labelled graph: concatenate with the nodes' features.



- ▶ Sensitive to the node ordering ! Also, expensive $O(|V|)$!
- ▶ Not applicable to graph with different sizes !

An encoder-decoder perspective

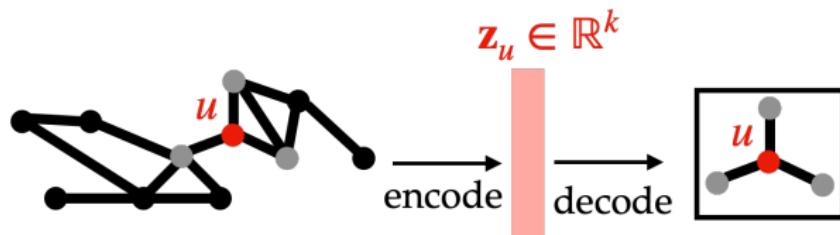
Notations

- ▶ We suppose we have one graph $G = (V, E)$, without features (so far).
- ▶ For each $u \in V$ we look for an embedding $\mathbf{z}_u \in \mathbb{R}^k$.

Principle

We look for a “good” encoder $E : V \rightarrow \mathbb{R}^k$ such that $E(u) = \mathbf{z}_u$.

- ▶ Ideally the embedding \mathbf{z}_u contains the neighbourhood informations of u .

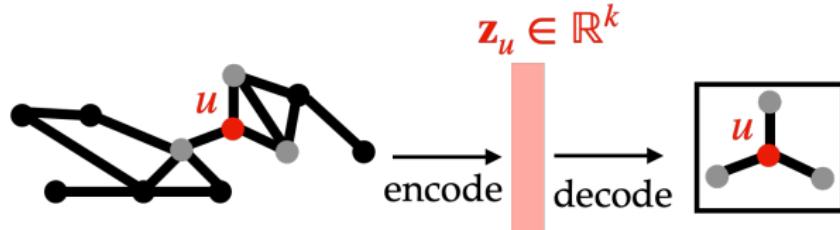


An encoder-decoder perspective

Principle

We look for a “good” encoder $E : V \rightarrow \mathbb{R}^k$ such that $E(u) = \mathbf{z}_u$.

- ▶ Ideally the embedding \mathbf{z}_u contains the neighbourhood informations of u .



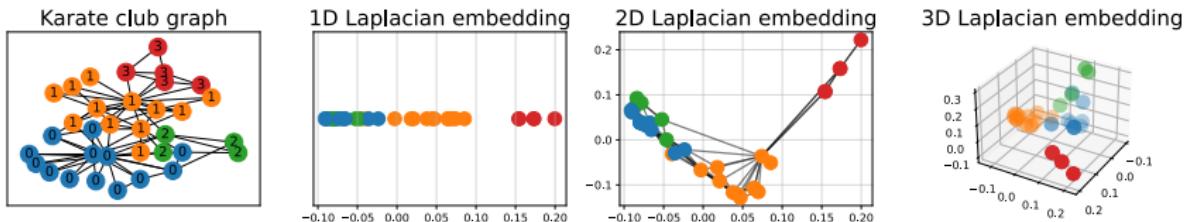
Encoding/decoding scheme

A lot of methods attempt to minimize

$$\mathcal{L} = \sum_{(u,v) \in \mathcal{D}} \ell(\text{similarity}(\mathbf{z}_u, \mathbf{z}_v), S[u, v])$$

- ▶ $\text{similarity}(\mathbf{z}_u, \mathbf{z}_v)$ how close are the embeddings.
- ▶ $S[u, v]$ how close are the nodes in the graph.
- ▶ $\ell : \mathbb{R} \times \mathbb{R} \rightarrow \mathbb{R}$ is a loss: how similar are the similarities.

Unsupervised node embeddings techniques



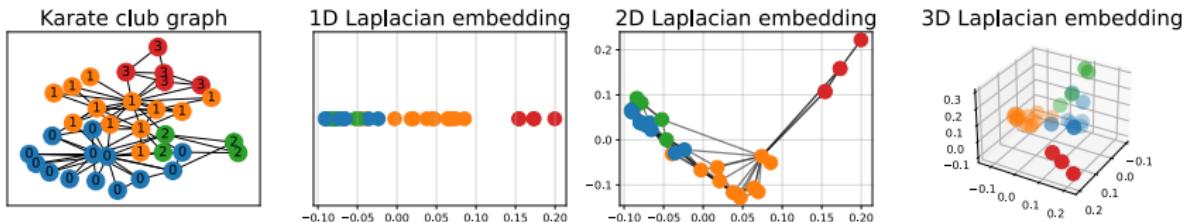
Inspiration from Laplacian eigenmaps Belkin and Niyogi 2003

- ▶ In the embedding space similarity($\mathbf{z}_u, \mathbf{z}_v$) = $\frac{1}{2} \|\mathbf{z}_u - \mathbf{z}_v\|_2^2$.
- ▶ When similiary is $\mathbf{S}[u_i, v_j] = A_{ij} / \sqrt{\text{degree}(u_i)} \sqrt{\text{degree}(u_j)}$, loss to minimize:

$$\frac{1}{2} \sum_{ij} \|\mathbf{z}_i - \mathbf{z}_j\|_2^2 \frac{A_{ij}}{\sqrt{\text{degree}(u_i)} \sqrt{\text{degree}(u_j)}} = \text{tr}(\mathbf{Z}^\top \tilde{\mathbf{L}} \mathbf{Z}).$$

- ▶ Normalized Laplacian $\tilde{\mathbf{L}} = \mathbf{D}^{-1/2} \mathbf{L} \mathbf{D}^{-1/2}$.
- ▶ Interpretation + permutation equivariance of the cost (on the board).

Unsupervised node embeddings techniques



Inspiration from Laplacian eigenmaps Belkin and Niyogi 2003

- ▶ In the embedding space similarity($\mathbf{z}_u, \mathbf{z}_v$) = $\frac{1}{2} \|\mathbf{z}_u - \mathbf{z}_v\|_2^2$.
- ▶ When similiary is $\mathbf{S}[u_i, v_j] = A_{ij} / \sqrt{\text{degree}(u_i)} \sqrt{\text{degree}(u_j)}$, loss to minimize:

$$\frac{1}{2} \sum_{ij} \|\mathbf{z}_i - \mathbf{z}_j\|_2^2 \frac{A_{ij}}{\sqrt{\text{degree}(u_i)} \sqrt{\text{degree}(u_j)}} = \text{tr}(\mathbf{Z}^\top \tilde{\mathbf{L}} \mathbf{Z}).$$

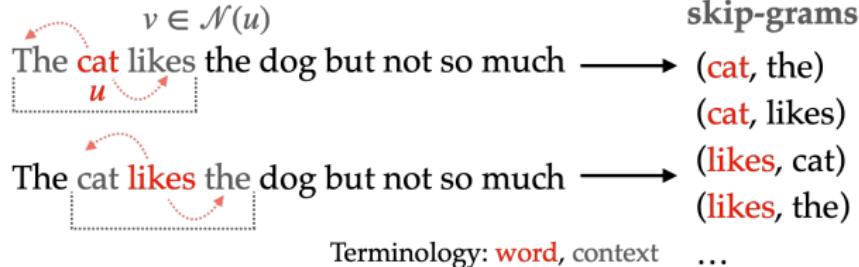
- ▶ Normalized Laplacian $\tilde{\mathbf{L}} = \mathbf{D}^{-1/2} \mathbf{L} \mathbf{D}^{-1/2}$.
- ▶ Interpretation + permutation equivariance of the cost (on the board).
- ▶ With the constraint $\mathbf{Z}^\top \mathbf{Z} = \mathbf{I}_d$ it recovers Laplacian eigenmaps.
- ▶ Sol. is the d eigenvectors associated to the d smallest eigenvalues of $\tilde{\mathbf{L}}$.

Unsupervised node embeddings techniques

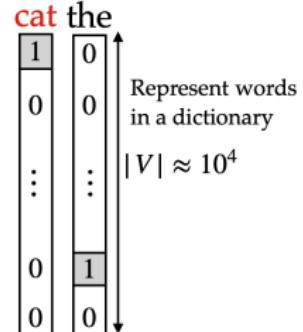
Skip-Gram and the Word2vec model (Mikolov et al. 2013)

The meaning of a word is its use in language (Wittgenstein).

- ▶ Objective: “similar” words are embedded into “similar” vectors.
- ▶ Goal: predict context words from each **input word**.
- ▶ We want to maximize $\mathbb{P}(\text{context}|\text{input word})$.



One hot encoding

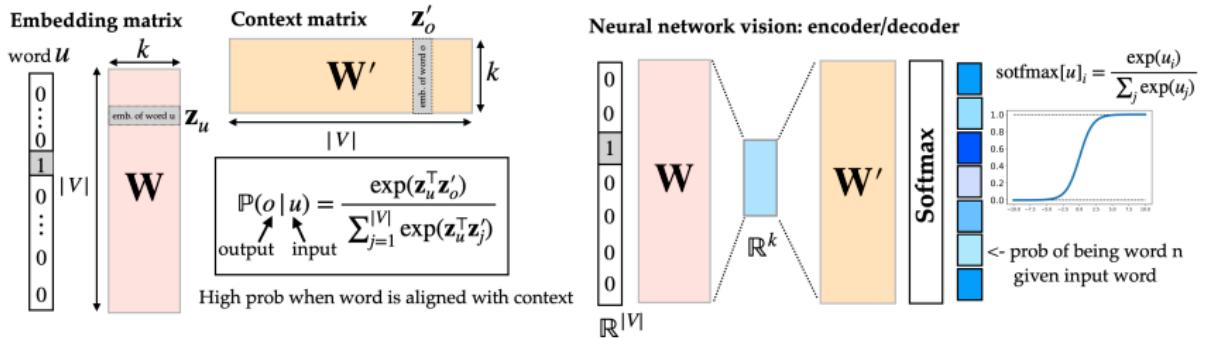


Unsupervised node embeddings techniques

Skip-Gram and the Word2vec model (Mikolov et al. 2013)

The meaning of a word is its use in language (Wittgenstein).

- ▶ Objective: “similar” words are embedded into “similar” vectors.
- ▶ Goal: predict context words from each **input word**.
- ▶ We want to maximize $\mathbb{P}(\text{context}|\text{input word})$.

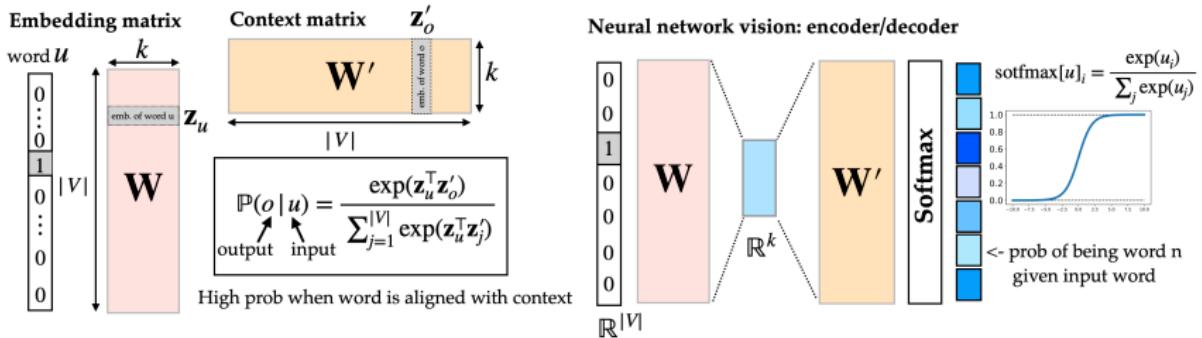


Unsupervised node embeddings techniques

Skip-Gram and the Word2vec model (Mikolov et al. 2013)

The meaning of a word is its use in language (Wittgenstein).

- ▶ Objective: “similar” words are embedded into “similar” vectors.
- ▶ Goal: predict context words from each **input word**.
- ▶ We want to maximize $\mathbb{P}(\text{context}|\text{input word})$.



- ▶ Dataset \mathcal{D} of input/output words (surrounding). Loss to minimize is:

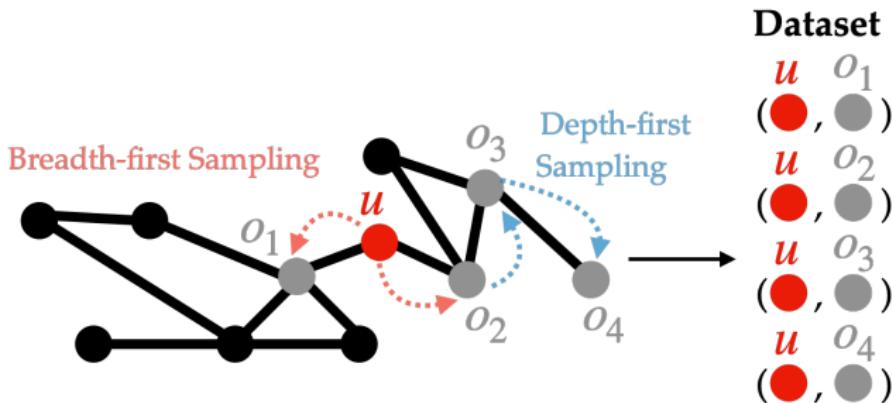
$$-\sum_{(u,o) \in \mathcal{D}} \log \mathbb{P}(o|u).$$

- ▶ But computing it in $\mathcal{O}(|V| \times |\{\text{words to embed}\}|)$: negative sampling.

Unsupervised node embeddings techniques

The node2vec model (Grover and Leskovec 2016)

- ▶ Similar as before: each node $u \in V$ is embedded as $\mathbf{z}_u \in \mathbb{R}^k$.
- ▶ Goal of the embedding: reflect the neighboring nodes of u .
- ▶ Sampling strategies based on random walks (BFS/DFS).



- ▶ With a dataset \mathcal{D} of input/output nodes. Loss to minimize:

$$\mathcal{L} = - \sum_{(u,o) \in \mathcal{D}} \log \frac{\exp(\mathbf{z}_u^\top \mathbf{z}_o)}{\sum_{w \in V} \exp(\mathbf{z}_u^\top \mathbf{z}_w)}$$

Unsupervised node embeddings techniques

The node2vec model (Grover and Leskovec 2016)

- ▶ Similar as before: each node $u \in V$ is embedded as $\mathbf{z}_u \in \mathbb{R}^k$.
- ▶ Goal of the embedding: reflect the neighboring nodes of u .
- ▶ Sampling strategies based on random walks (BFS/DFS).

Negative sampling (NS)

- ▶ Loss is too expensive to compute $\mathcal{O}(|V|^2)$.
- ▶ NS: introduce negative data samples.
- ▶ Goal: distinguish between neighboring points of a target node u and random nodes drawn from a noise distribution using logistic regression.
- ▶ New loss (explanations on the board) (Goldberg and Levy 2014):

$$\mathcal{L} = - \left(\sum_{(u_+, o_+) \in \mathcal{D}_+} \log \sigma(\mathbf{z}_u^\top \mathbf{z}_o) + \sum_{(u_-, o_-) \in \mathcal{D}_-} \log \sigma(-\mathbf{z}_u^\top \mathbf{z}_o) \right)$$

with sigmoid function $\sigma(x) = \frac{1}{1 + \exp(-x)}$.

Unsupervised node embeddings techniques

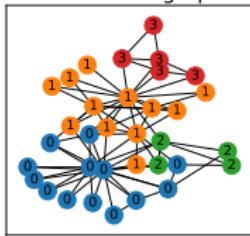
The node2vec model (Grover and Leskovec 2016)

- ▶ Similar as before: each node $u \in V$ is embedded as $\mathbf{z}_u \in \mathbb{R}^k$.
- ▶ Goal of the embedding: reflect the neighboring nodes of u .
- ▶ Sampling strategies based on random walks (BFS/DFS).

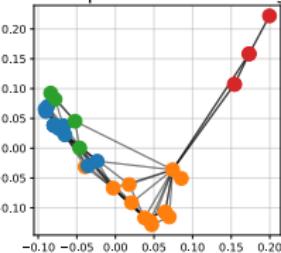
Negative sampling (NS)

- ▶ Goal: distinguish between neighboring points of a target node u and random draws from a noise distribution using logistic regression.

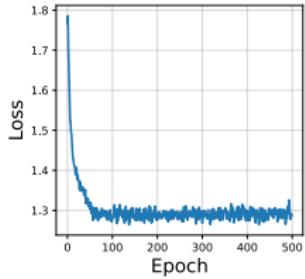
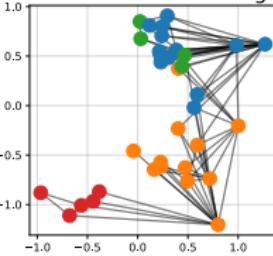
Karate club graph



2D Laplacian embedding



node2vec embedding



Unsupervised node embeddings techniques

Limitations of previous embeddings techniques

- ▶ The previous embeddings are called **shallow**: encoder function $E : V \rightarrow \mathbb{R}^k$ is simply an embedding lookup based on the node ID.

$$E(u) = \mathbf{Z}[:, u] = \mathbf{z}_u .$$

Unsupervised node embeddings techniques

Limitations of previous embeddings techniques

- ▶ The previous embeddings are called **shallow**: encoder function $E : V \rightarrow \mathbb{R}^k$ is simply an embedding lookup based on the node ID.

$$E(u) = \mathbf{Z}[:, u] = \mathbf{z}_u .$$

- ▶ Lack of parameter sharing between nodes in the encoder.
- ▶ Do not leverage node features !
- ▶ Inherently **transductive**: these methods can only generate embeddings for nodes that were present during the training phase.
- ▶ **If new nodes must retrain everything.**

References I

-  Belkin, Mikhail and Partha Niyogi (2003). "Laplacian eigenmaps for dimensionality reduction and data representation". In: *Neural computation* 15.6, pp. 1373–1396.
-  Bruna, Joan et al. (2013). "Spectral networks and locally connected networks on graphs". In: *arXiv preprint arXiv:1312.6203*.
-  Defferrard, Michaël, Xavier Bresson, and Pierre Vandergheynst (2017). *Convolutional Neural Networks on Graphs with Fast Localized Spectral Filtering*. *arXiv: 1606.09375 [cs.LG]*.
-  Goldberg, Yoav and Omer Levy (2014). "word2vec Explained: deriving Mikolov et al.'s negative-sampling word-embedding method". In: *arXiv preprint arXiv:1402.3722*.
-  Grinsztajn, Léo, Edouard Oyallon, and Gaël Varoquaux (2022). "Why do tree-based models still outperform deep learning on typical tabular data?" In: *Advances in Neural Information Processing Systems* 35, pp. 507–520.
-  Grover, Aditya and Jure Leskovec (2016). "node2vec: Scalable feature learning for networks". In: *Proceedings of the 22nd ACM SIGKDD international conference on Knowledge discovery and data mining*, pp. 855–864.

References II

-  Hamilton, William L (2020). *Graph representation learning*. Morgan & Claypool Publishers.
-  Hamilton, William L., Rex Ying, and Jure Leskovec (2018). *Inductive Representation Learning on Large Graphs*. arXiv: 1706.02216 [cs.SI].
-  Keriven, Nicolas (2022). "Not too little, not too much: a theoretical analysis of graph (over) smoothing". In: *Advances in Neural Information Processing Systems 35*, pp. 2268–2281.
-  Kipf, Thomas N and Max Welling (2016). "Semi-supervised classification with graph convolutional networks". In: *arXiv preprint arXiv:1609.02907*.
-  LeCun, Yann et al. (1998). "Gradient-based learning applied to document recognition". In: *Proceedings of the IEEE 86.11*, pp. 2278–2324.
-  Luan, Sitao et al. (2022). "Revisiting heterophily for graph neural networks". In: *Advances in neural information processing systems 35*, pp. 1362–1375.
-  Maron, Haggai et al. (2018). "Invariant and equivariant graph networks". In: *arXiv preprint arXiv:1812.09902*.
-  Mikolov, Tomas et al. (2013). "Distributed representations of words and phrases and their compositionality". In: *Advances in neural information processing systems 26*.

References III

-  Rosenblatt, Frank (1958). "The perceptron: a probabilistic model for information storage and organization in the brain.". In: *Psychological review* 65.6, p. 386.
-  Scarselli, Franco et al. (2008). "The graph neural network model". In: *IEEE transactions on neural networks* 20.1, pp. 61–80.
-  Vaswani, Ashish et al. (2023). *Attention Is All You Need*. arXiv: 1706.03762 [cs.CL].
-  Velivcković, Petar et al. (2017). "Graph attention networks". In: *arXiv preprint arXiv:1710.10903*.
-  Wagstaff, Edward et al. (2019). "On the limitations of representing functions on sets". In: *International Conference on Machine Learning*. PMLR, pp. 6487–6494.
-  Weiler, Maurice et al. (2023). *Equivariant and Coordinate Independent Convolutional Networks. A Gauge Field Theory of Neural Networks*. URL: https://maurice-weiler.gitlab.io/cnn_book/EquivariantAndCoordinateIndependentCNNs.pdf.

References IV

-  Wu, Zonghan et al. (Jan. 2021). "A Comprehensive Survey on Graph Neural Networks". In: *IEEE Transactions on Neural Networks and Learning Systems* 32.1, pp. 4–24. DOI: 10.1109/tnnls.2020.2978386. URL: <https://doi.org/10.1109%2Ftnnls.2020.2978386>.
-  Xu, Keyulu et al. (2019). *How Powerful are Graph Neural Networks?* arXiv: 1810.00826 [cs.LG].
-  Ying, Zhitao et al. (2018). "Hierarchical graph representation learning with differentiable pooling". In: *Advances in neural information processing systems* 31.
-  Zaheer, Manzil et al. (2018). *Deep Sets*. arXiv: 1703.06114 [cs.LG].
-  Zeiler, Matthew D and Rob Fergus (2014). "Visualizing and understanding convolutional networks". In: *Computer Vision–ECCV 2014: 13th European Conference, Zurich, Switzerland, September 6–12, 2014, Proceedings, Part I* 13. Springer, pp. 818–833.
-  Zhang, Ziwei, Peng Cui, and Wenwu Zhu (2020). *Deep Learning on Graphs: A Survey*. arXiv: 1812.04202 [cs.LG].