

Machine learning for graphs and with graphs

Course 1: introduction

Titouan Vayer & Pierre Borgnat

email: titouan.vayer@inria.fr, pierre.borgnat@ens-lyon.fr

September 4, 2024



Machine learning for graphs and with graphs

From theory ...

1. Basics of machine learning
2. The graph framework
3. Community detection/ graph clustering
4. Graph signal processing
5. Kernels for graphs
6. Graph neural networks
7. Optimal transport for graphs
8. Learning graphs from (unstructured) data

... to practice

We will use Python and various libraries



Ressources:

<https://tvayer.github.io/courses/coursegraph.html>

Some references for machine learning

Shai Shalev-Shwartz and Shai Ben-David (2014).

Understanding Machine Learning - From Theory to Algorithms. Cambridge University Press

Francis Bach (2022). *Learning Theory from First Principles.*

Trevor Hastie, Robert Tibshirani, and Jerome Friedman (2001). *The Elements of Statistical Learning.* Springer New York Inc.

Evaluation

- ▶ 50 % oral presentation on a selected research article.
- ▶ 50 % code associated to the article applied on real data.
- ▶ Bonus.

Python installations

- ▶ The practical sessions of the course will require to run jupyter notebooks.
- ▶ We recommend that you install python through the Anaconda distribution (python 3.7, 3.8 or 3.9 is preferable) available at <https://www.anaconda.com/products/distribution>

You should check that you are able to create and open a jupyter notebook, and inside, run the following imports:

```
1 import matplotlib
2 import numpy
3 import sklearn
4 import pytorch
5 import pandas
6 import scipy
```

If any of these packages is missing, it can be installed with ‘conda install numpy’, the command being run in a terminal or in Anaconda prompt for Windows user.

Basics of machine learning

What is machine learning ?

Data in machine learning

From training data to prediction

Loss functions

Empirical risk minimization

Train/validation/test

Model selection and validation

Split your dataset !

The problems with structured data

Motivating examples

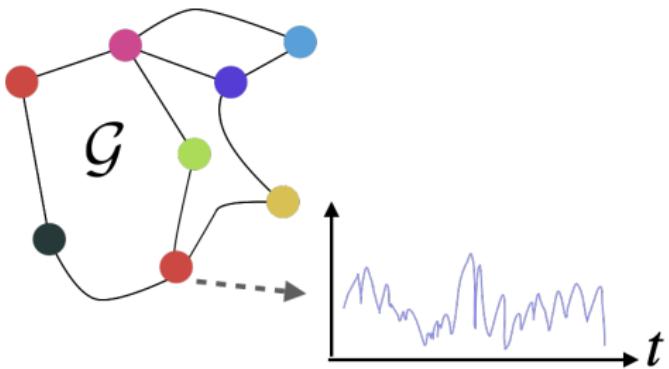
A primer on graph theory

Why “classical ML” struggles with structured data

What is machine learning ?

Some applications

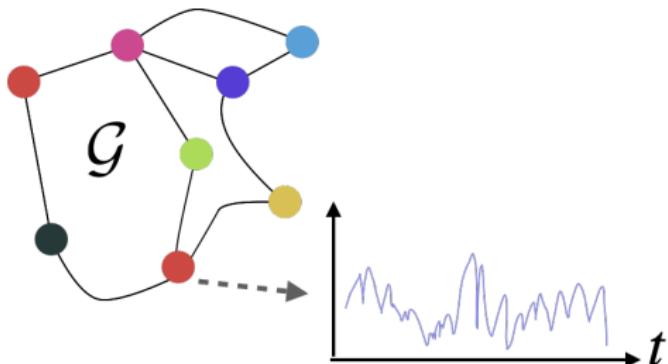
1. Energy networks, disease propagation



What is machine learning ?

Some applications

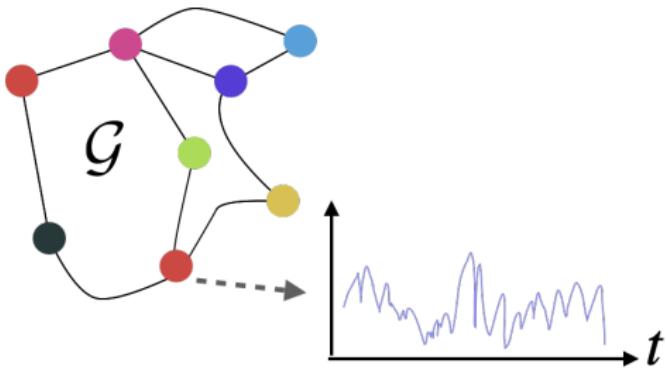
1. Energy networks, disease propagation
2. Image analysis (medical application, web)



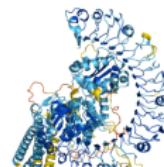
What is machine learning ?

Some applications

1. Energy networks, disease propagation
2. Image analysis (medical application, web)
3. Protein folding Jumper et al. 2021



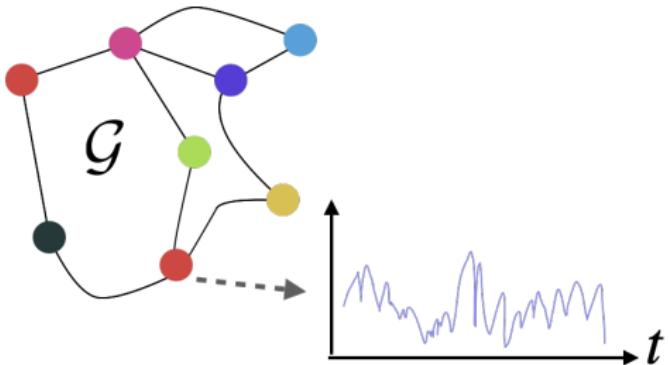
AAATGCG.... ----->



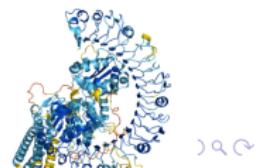
What is machine learning ?

Some applications

1. Energy networks, disease propagation
2. Image analysis (medical application, web)
3. Protein folding Jumper et al. 2021
4. Generative models <https://stablediffusionweb.com/>



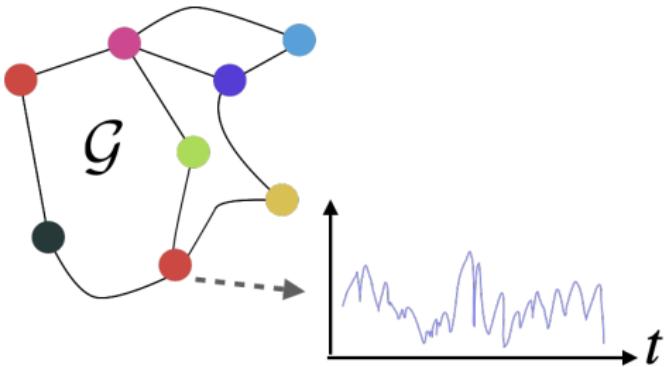
AAATGCG.... ----->



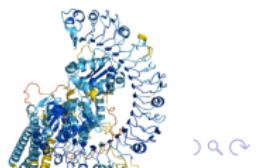
What is machine learning ?

Some applications

1. Energy networks, disease propagation
2. Image analysis (medical application, web)
3. Protein folding Jumper et al. 2021
4. Generative models <https://stablediffusionweb.com/>
5. Natural language processing
<https://chat.openai.com/chat>



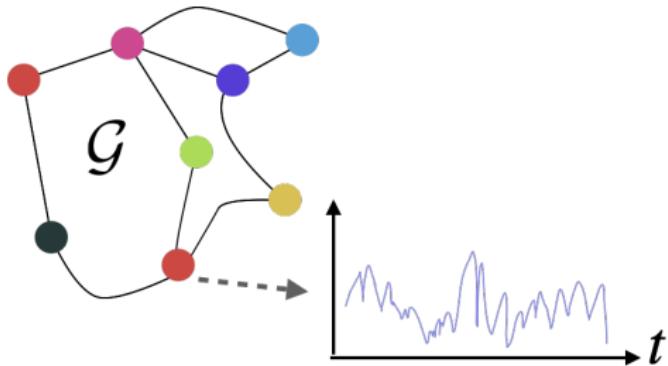
AAATGCG.... →



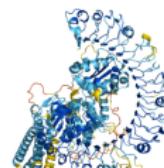
What is machine learning ?

Some applications

1. Energy networks, disease propagation
2. Image analysis (medical application, web)
3. Protein folding Jumper et al. 2021
4. Generative models <https://stablediffusionweb.com/>
5. Natural language processing
<https://chat.openai.com/chat>
6. For art <https://www.youtube.com/watch?v=MwtVkJx3RA>



AAATGCG.... ----->

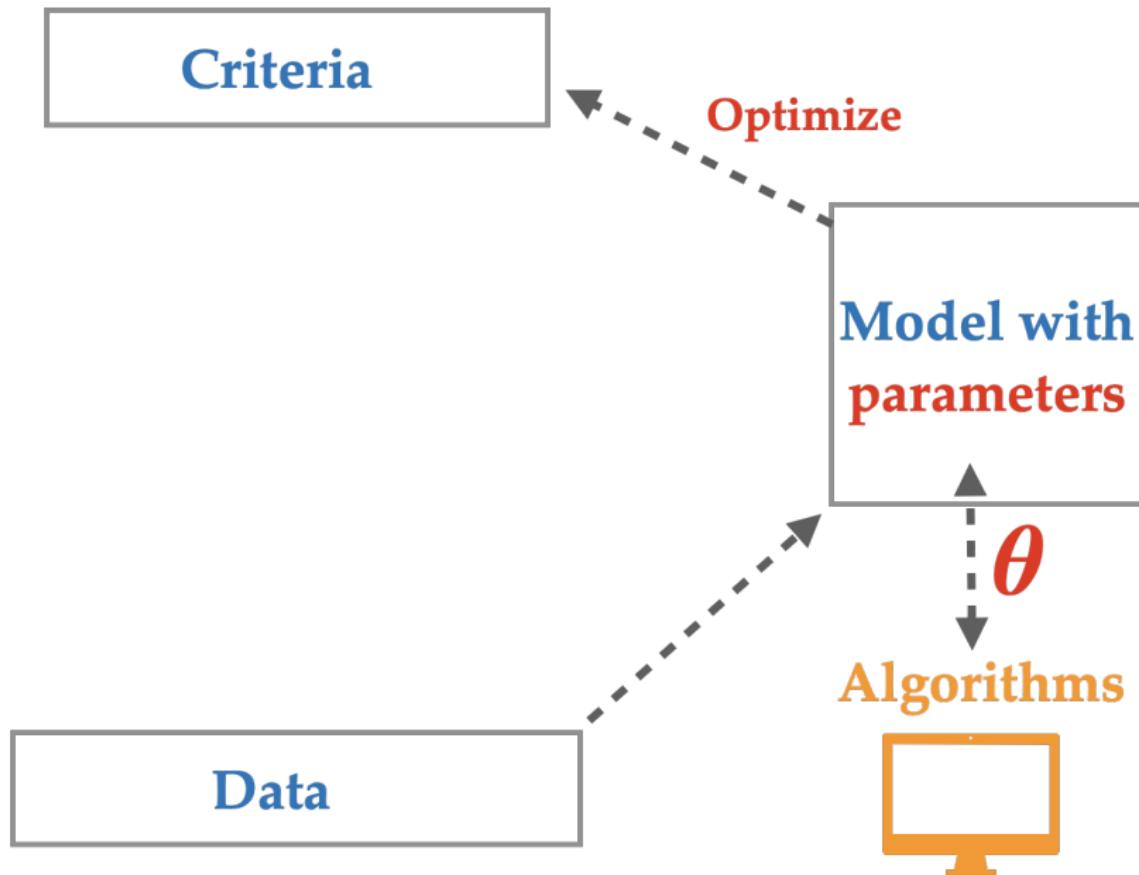


What is machine learning ?

Criteria

Data

What is machine learning ?



What is machine learning ?

The objective of machine learning

Teach a machine to process automatically some data in order to solve a given problem.

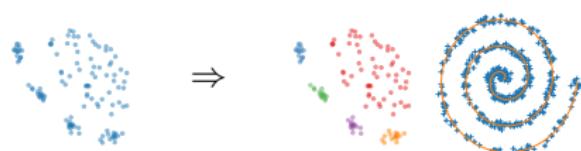
What is machine learning ?

The objective of machine learning

Teach a machine to process automatically some data in order to solve a given problem.

Unsupervised learning: understanding the data

- ▶ Clustering & probability density estimation
- ▶ Dimensionality reduction



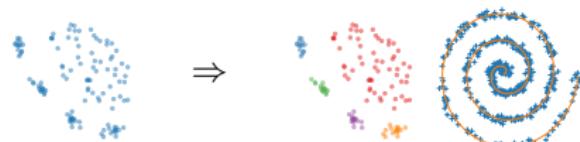
What is machine learning ?

The objective of machine learning

Teach a machine to process automatically some data in order to solve a given problem.

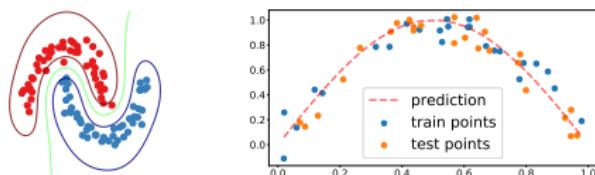
Unsupervised learning: understanding the data

- ▶ Clustering & probability density estimation
- ▶ Dimensionality reduction



Supervised learning: learning to predict

- ▶ Classification: classify points according to some labels
- ▶ Regression: predict real (vector) values



Some images and slides have been obtained by the courtesy of Rémi Flamary

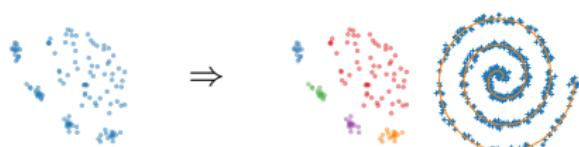
What is machine learning ?

The objective of machine learning

Teach a machine to process automatically some data in order to solve a given problem.

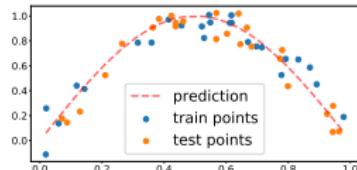
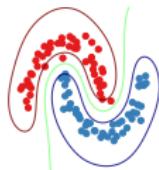
Unsupervised learning: understanding the data

- ▶ Clustering & probability density estimation
- ▶ Dimensionality reduction

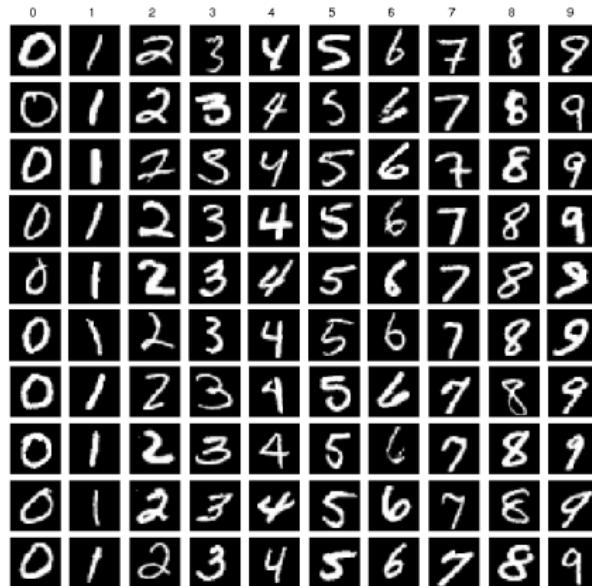


Supervised learning: learning to predict

- ▶ Classification: classify points according to some labels
- ▶ Regression: predict real (vector) values



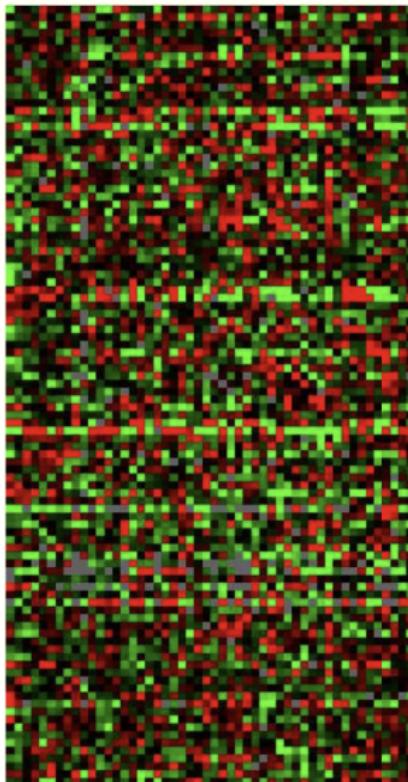
What is machine learning ?



Supervised classification examples

- ▶ e.g. to identify the numbers on images from a 16×16 gray level image (image classification)
- ▶ SPAM, fraud detection, disease classification ...

What is machine learning ?



Clustering example

- ▶ Analyse n sequences (individuals) of d genetical responses
- ▶ Groups of similar samples ? Gene with similar expressions ?

Plan

What is machine learning ?

Data in machine learning

From training data to prediction

Loss functions

Empirical risk minimization

Train/validation/test

Model selection and validation

Split your dataset !

The problems with structured data

Motivating examples

A primer on graph theory

Why “classical ML” struggles with structured data

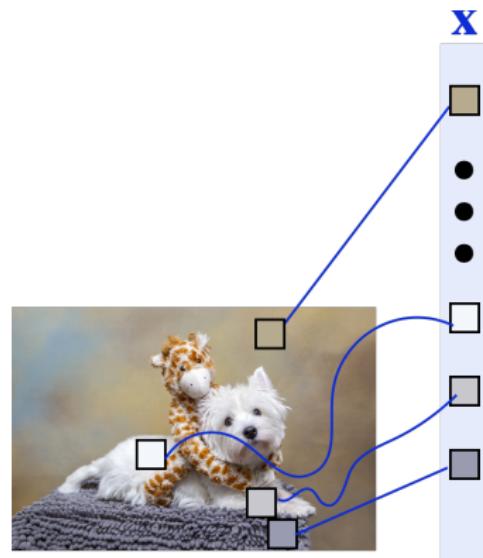
Store a data point

Vectorial representation

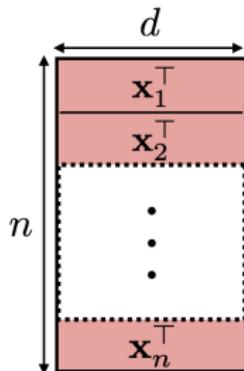
One “sample”, “data point”, “individual”:

$$\mathbf{x} = (x_1, \dots, x_d)^\top \in \mathbb{R}^d$$

- ▶ d is the dimension, x_i is the i th information i of \mathbf{x}
- ▶ Can describe information about an individual
- ▶ For an image \mathbf{x} : each pixel of an image
- ▶ Descriptors of a cell, word embedding ...



Unsupervised dataset

$$\mathbf{X} = \begin{bmatrix} \mathbf{x}_1^\top \\ \mathbf{x}_2^\top \\ \vdots \\ \mathbf{x}_n^\top \end{bmatrix} = \begin{bmatrix} x_{11} & x_{12} & \dots & x_{1d} \\ \vdots & \vdots & \vdots & \vdots \\ x_{n1} & x_{n2} & \dots & x_{nd} \end{bmatrix}$$


Unsupervised learning

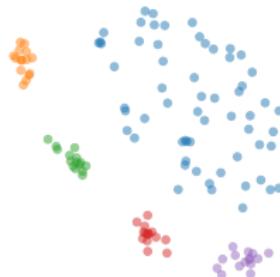
- ▶ The dataset contains the samples $(\mathbf{x}_i)_{i=1}^n$ where n is the number of samples of size d .
- ▶ d and n define the dimensionality of the learning problem.
- ▶ Data stored as a matrix $\mathbf{X} \in \mathbb{R}^{n \times d}$ that contains the training samples as rows.

Supervised dataset

Samples + labels:

$$\mathbf{X} = \begin{bmatrix} \mathbf{x}_1^\top \\ \mathbf{x}_2^\top \\ \vdots \\ \mathbf{x}_n^\top \end{bmatrix} \quad \mathbf{y} = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}$$

Classification



Regression



Supervised learning

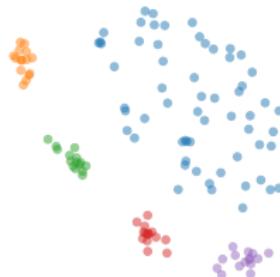
- ▶ The dataset contains the samples $(\mathbf{x}_i, y_i)_{i=1}^n$ where \mathbf{x}_i is the feature sample and $y_i \in \mathcal{Y}$ its label.
- ▶ The values to predict (label) can be concatenated in a vector $\mathbf{y} \in \mathcal{Y}^n$

Supervised dataset

Samples + labels:

$$\mathbf{X} = \begin{bmatrix} \mathbf{x}_1^\top \\ \mathbf{x}_2^\top \\ \vdots \\ \mathbf{x}_n^\top \end{bmatrix} \quad \mathbf{y} = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}$$

Classification



Regression



Supervised learning

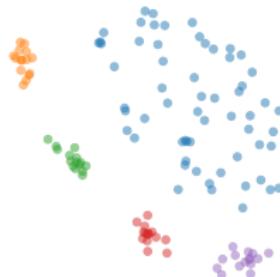
- ▶ The dataset contains the samples $(\mathbf{x}_i, y_i)_{i=1}^n$ where \mathbf{x}_i is the feature sample and $y_i \in \mathcal{Y}$ its label.
- ▶ The values to predict (label) can be concatenated in a vector $\mathbf{y} \in \mathcal{Y}^n$
- ▶ Prediction space \mathcal{Y} can be:
 - ▶ $\mathcal{Y} = \{-1, 1\}$ or $\mathcal{Y} = \{1, \dots, K\}$ for classification problems.
 - ▶ $\mathcal{Y} = \mathbb{R}$ for regression problems (\mathbb{R}^p for multi-output regression).

Supervised dataset

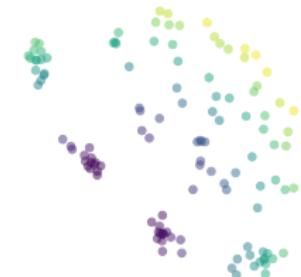
Samples + labels:

$$\mathbf{X} = \begin{bmatrix} \mathbf{x}_1^\top \\ \mathbf{x}_2^\top \\ \vdots \\ \mathbf{x}_n^\top \end{bmatrix} \quad \mathbf{y} = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}$$

Classification



Regression



Supervised learning

- ▶ The dataset contains the samples $(\mathbf{x}_i, y_i)_{i=1}^n$ where \mathbf{x}_i is the feature sample and $y_i \in \mathcal{Y}$ its label.
- ▶ The values to predict (label) can be concatenated in a vector $\mathbf{y} \in \mathcal{Y}^n$
- ▶ Prediction space \mathcal{Y} can be:
 - ▶ $\mathcal{Y} = \{-1, 1\}$ or $\mathcal{Y} = \{1, \dots, K\}$ for classification problems.
 - ▶ $\mathcal{Y} = \mathbb{R}$ for regression problems (\mathbb{R}^p for multi-output regression).
- ▶ Semi-supervised learning: few labeled points are available, but a large number of unlabeled points are given.

Plan

What is machine learning ?

Data in machine learning

From training data to prediction

Loss functions

Empirical risk minimization

Train/validation/test

Model selection and validation

Split your dataset !

The problems with structured data

Motivating examples

A primer on graph theory

Why “classical ML” struggles with structured data

From training data to prediction

Training data

- ▶ We have access to n samples $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n) \sim p$
- ▶ $p \in \mathcal{P}(\mathcal{X} \times \mathcal{Y})$ is the data distribution
- ▶ p is unknown ! We only have access to samples.

From training data to prediction

Training data

- ▶ We have access to n samples $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n) \sim p$
- ▶ $p \in \mathcal{P}(\mathcal{X} \times \mathcal{Y})$ is the data distribution
- ▶ p is unknown ! We only have access to samples.
- ▶ For unsupervised problem we only have $\mathbf{x}_1, \dots, \mathbf{x}_n \sim p$ and $p \in \mathcal{P}(\mathcal{X})$

From training data to prediction

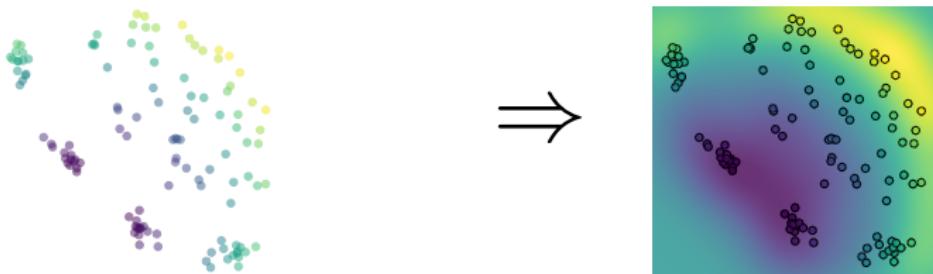
Training data

- ▶ We have access to n samples $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n) \sim p$
- ▶ $p \in \mathcal{P}(\mathcal{X} \times \mathcal{Y})$ is the data distribution
- ▶ p is unknown ! We only have access to samples.
- ▶ For unsupervised problem we only have $\mathbf{x}_1, \dots, \mathbf{x}_n \sim p$ and $p \in \mathcal{P}(\mathcal{X})$

Objective

- ▶ We have a task to solve: classification, regression, clustering ...
- ▶ Most ML problems formulate as **finding some function f that “best” solves our task**
- ▶ f is called **an hypothesis** and is **implemented by a computer**
- ▶ Most of the time f depends on some parameter $\theta \in \Theta$

Regression

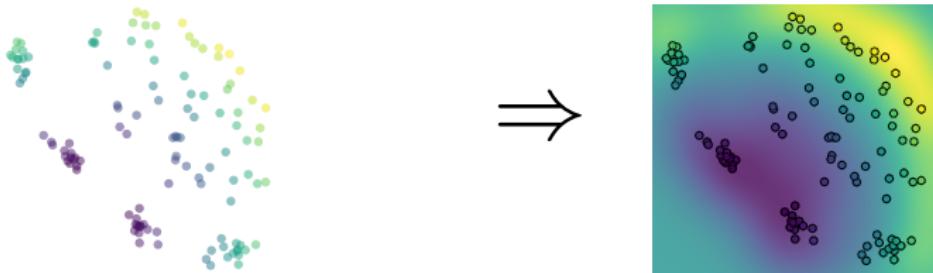


Objective

$$(\mathbf{x}_i, y_i)_{i=1}^n \Rightarrow f : \mathbb{R}^d \rightarrow \mathbb{R}$$

- ▶ Train a function $f(\mathbf{x}) = y \in \mathcal{Y}$ predicting a continuous value ($\mathcal{Y} = \mathbb{R}$).
- ▶ Can be extended to multi-value prediction ($\mathcal{Y} = \mathbb{R}^p$).

Regression



Objective

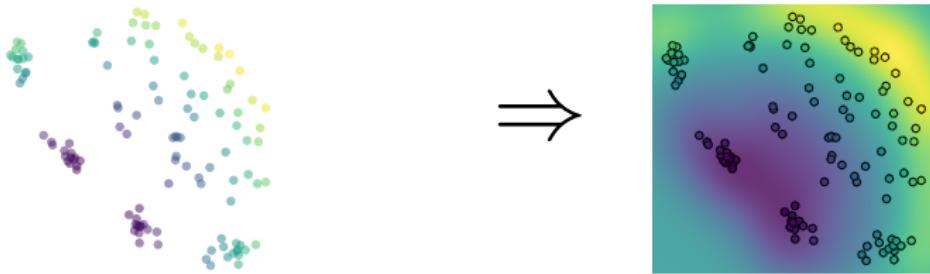
$$(\mathbf{x}_i, y_i)_{i=1}^n \Rightarrow f : \mathbb{R}^d \rightarrow \mathbb{R}$$

- ▶ Train a function $f(\mathbf{x}) = y \in \mathcal{Y}$ predicting a continuous value ($\mathcal{Y} = \mathbb{R}$).
- ▶ Can be extended to multi-value prediction ($\mathcal{Y} = \mathbb{R}^p$).

Hyperparameters

- ▶ Type of function (linear, kernel, neural network).
- ▶ Performance measure.
- ▶ Regularization.

Regression



Objective

$$(\mathbf{x}_i, y_i)_{i=1}^n \Rightarrow f : \mathbb{R}^d \rightarrow \mathbb{R}$$

- ▶ Train a function $f(\mathbf{x}) = y \in \mathcal{Y}$ predicting a continuous value ($\mathcal{Y} = \mathbb{R}$).
- ▶ Can be extended to multi-value prediction ($\mathcal{Y} = \mathbb{R}^p$).

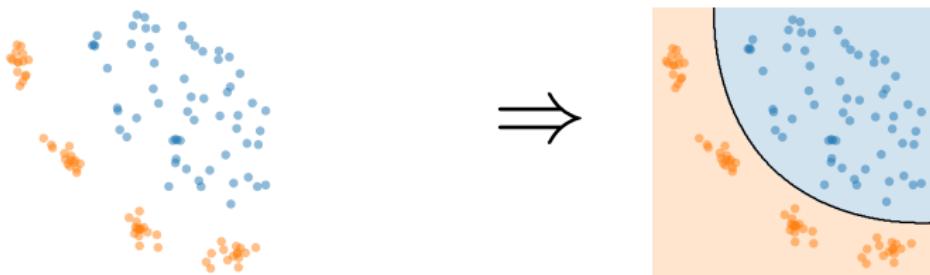
Hyperparameters

- ▶ Type of function (linear, kernel, neural network).
- ▶ Performance measure.
- ▶ Regularization.

Methods

- ▶ Least Square (LS).
- ▶ Ridge regression, Lasso.
- ▶ Kernel regression.
- ▶ Deep learning.

Binary classification

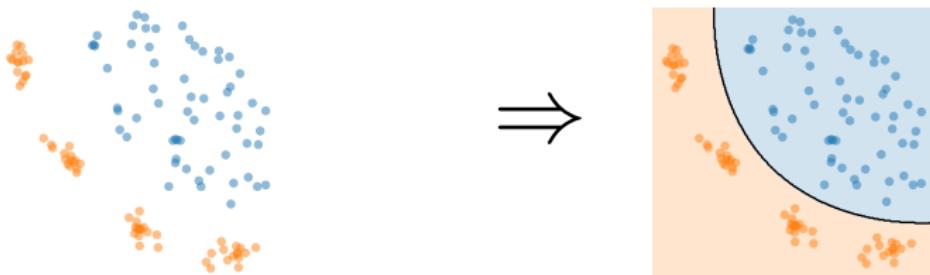


Objective

$$(\mathbf{x}_i, y_i)_{i=1}^n \Rightarrow f : \mathbb{R}^d \rightarrow \{-1, 1\}$$

- ▶ Train a function $f(\mathbf{x}) = y \in \mathcal{Y}$ predicting a binary value ($\mathcal{Y} = \{-1, 1\}$).
- ▶ $f(\mathbf{x}) = 0$ defines the boundary on the partition of the feature space.

Binary classification



Objective

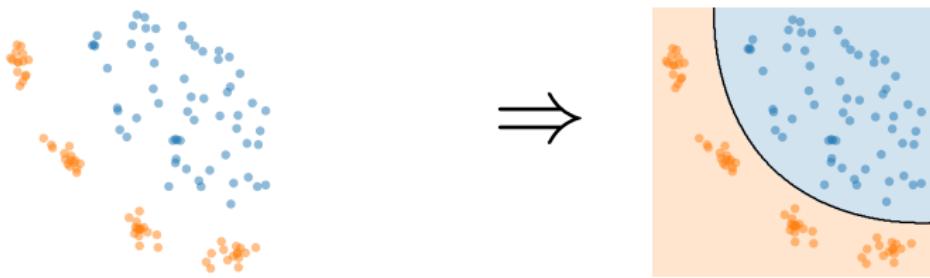
$$(\mathbf{x}_i, y_i)_{i=1}^n \Rightarrow f : \mathbb{R}^d \rightarrow \{-1, 1\}$$

- ▶ Train a function $f(\mathbf{x}) = y \in \mathcal{Y}$ predicting a binary value ($\mathcal{Y} = \{-1, 1\}$).
- ▶ $f(\mathbf{x}) = 0$ defines the boundary on the partition of the feature space.

Hyperparameters

- ▶ Type of function (linear, kernel, neural network).
- ▶ Performance measure.
- ▶ Regularization.

Binary classification



Objective

$$(\mathbf{x}_i, y_i)_{i=1}^n \Rightarrow f : \mathbb{R}^d \rightarrow \{-1, 1\}$$

- ▶ Train a function $f(\mathbf{x}) = y \in \mathcal{Y}$ predicting a binary value ($\mathcal{Y} = \{-1, 1\}$).
- ▶ $f(\mathbf{x}) = 0$ defines the boundary on the partition of the feature space.

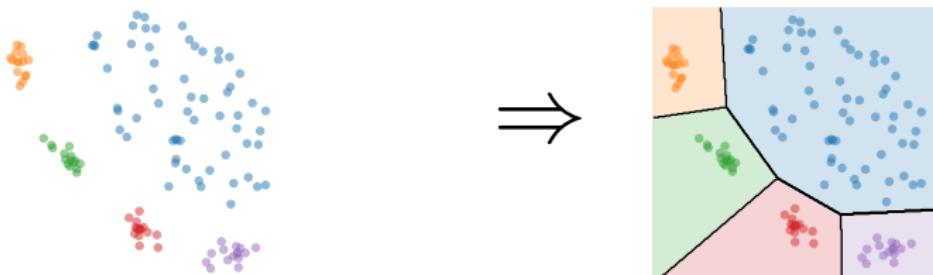
Hyperparameters

- ▶ Type of function (linear, kernel, neural network).
- ▶ Performance measure.
- ▶ Regularization.

Methods

- ▶ Bayesian classifier (LDA, QDA)
- ▶ Linear and kernel discrimination
- ▶ Decision trees, random forests.
- ▶ Deep learning.

Multiclass classification



Objective

$$(\mathbf{x}_i, y_i)_{i=1}^n \Rightarrow f : \mathbb{R}^d \rightarrow \{1, \dots, K\}$$

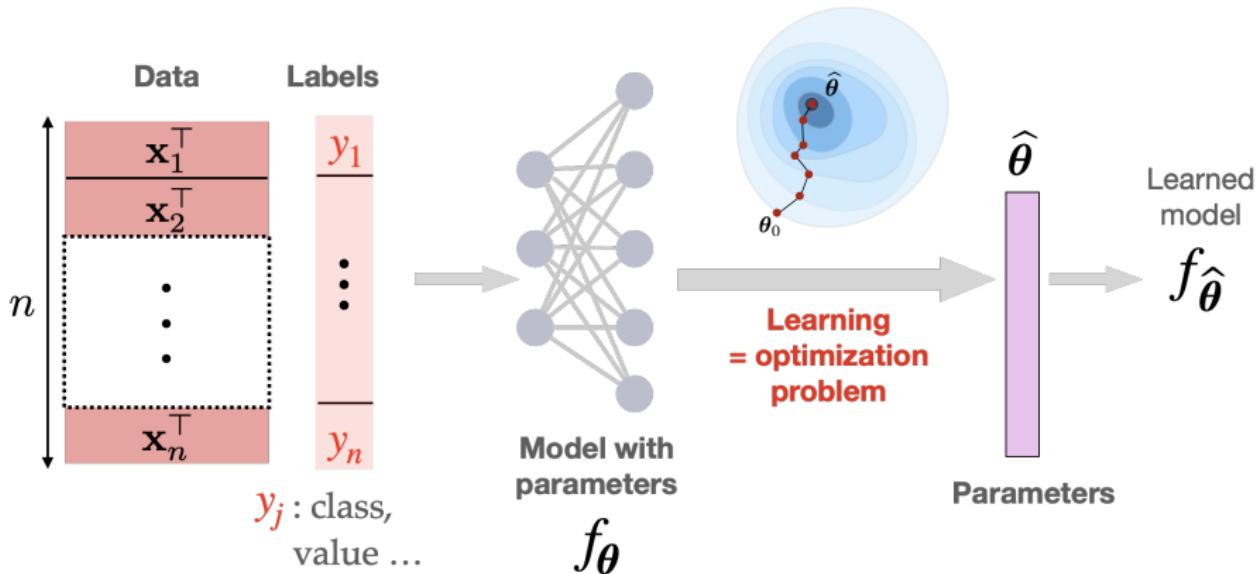
- ▶ Train a function $f(\mathbf{x}) = y \in \mathcal{Y}$ predicting an integer value ($\mathcal{Y} = \{1, \dots, K\}$).
- ▶ In practice K continuous score functions f_k are estimated and the prediction is

$$f(\mathbf{x}) = \arg \max_k f_k(\mathbf{x})$$

- ▶ Softmax can be used instead of argmax to get probability estimates.

The big picture of (parametrized) ML

But how to find this function ?



The goal in the **learning step** will be to find the parameters $\hat{\theta}$ (hence the function) that **minimizes a measure of error on the data**

Loss functions

Supervised case

A loss function is $\ell : \mathcal{Y} \times \mathcal{Y} \rightarrow \mathbb{R}$ so that:

$$\ell(\text{true value}, \text{predicted value}) = \text{ how good is my prediction}$$

Regression problems

Loss functions

Supervised case

A loss function is $\ell : \mathcal{Y} \times \mathcal{Y} \rightarrow \mathbb{R}$ so that:

$$\ell(\text{true value}, \text{predicted value}) = \text{ how good is my prediction}$$

Regression problems

- ▶ E.g. $y_i \in \mathbb{R}$ $\ell(y_i, f(\mathbf{x}_i)) = (y_i - f(\mathbf{x}_i))^2$ (square loss)

Loss functions

Supervised case

A loss function is $\ell : \mathcal{Y} \times \mathcal{Y} \rightarrow \mathbb{R}$ so that:

$$\ell(\text{true value}, \text{predicted value}) = \text{ how good is my prediction}$$

Regression problems

- ▶ E.g. $y_i \in \mathbb{R}$ $\ell(y_i, f(\mathbf{x}_i)) = (y_i - f(\mathbf{x}_i))^2$ (square loss)
- ▶ E.g. $\mathbf{y}_i \in \mathbb{R}^p$ $\ell(\mathbf{y}_i, f(\mathbf{x}_i)) = \|\mathbf{y}_i - f(\mathbf{x}_i)\|_2^2$ (square loss)

Loss functions

Supervised case

A loss function is $\ell : \mathcal{Y} \times \mathcal{Y} \rightarrow \mathbb{R}$ so that:

$$\ell(\text{true value}, \text{predicted value}) = \text{ how good is my prediction}$$

Regression problems

- ▶ E.g. $y_i \in \mathbb{R}$ $\ell(y_i, f(\mathbf{x}_i)) = (y_i - f(\mathbf{x}_i))^2$ (square loss)
- ▶ E.g. $\mathbf{y}_i \in \mathbb{R}^p$ $\ell(\mathbf{y}_i, f(\mathbf{x}_i)) = \|\mathbf{y}_i - f(\mathbf{x}_i)\|_2^2$ (square loss)
- ▶ E.g. $\mathbf{y}_i \in \mathbb{R}^p$ $\ell(\mathbf{y}_i, f(\mathbf{x}_i)) = \|\mathbf{y}_i - f(\mathbf{x}_i)\|_q^q$ (ℓ_q loss)

Loss functions

Supervised case

A loss function is $\ell : \mathcal{Y} \times \mathcal{Y} \rightarrow \mathbb{R}$ so that:

$$\ell(\text{true value}, \text{predicted value}) = \text{ how good is my prediction}$$

Regression problems

- ▶ E.g. $y_i \in \mathbb{R}$ $\ell(y_i, f(\mathbf{x}_i)) = (y_i - f(\mathbf{x}_i))^2$ (square loss)
- ▶ E.g. $\mathbf{y}_i \in \mathbb{R}^p$ $\ell(\mathbf{y}_i, f(\mathbf{x}_i)) = \|\mathbf{y}_i - f(\mathbf{x}_i)\|_2^2$ (square loss)
- ▶ E.g. $\mathbf{y}_i \in \mathbb{R}^p$ $\ell(\mathbf{y}_i, f(\mathbf{x}_i)) = \|\mathbf{y}_i - f(\mathbf{x}_i)\|_q^q$ (ℓ_q loss)

Classification problems

- ▶ E.g. $y_i \in \{-1, 1\}$ $\ell(y_i, f(\mathbf{x}_i)) = \mathbf{1}_{y_i \neq f(\mathbf{x}_i)}$ (0/1 loss)

Empirical risk minimization

Train by minimizing the train error

To find f the idea is to **minimize the averaged error** on the training samples:

$$\min_f \frac{1}{n} \sum_{i=1}^n \ell(y_i, f(\mathbf{x}_i)) \quad (\text{ERM})$$

Empirical risk minimization

Train by minimizing the train error

To find f the idea is to **minimize the averaged error** on the training samples:

$$\min_f \frac{1}{n} \sum_{i=1}^n \ell(y_i, f(\mathbf{x}_i)) \quad (\text{ERM})$$

- ▶ It is called **empirical risk minimization (ERM)**
- ▶ Given the loss, finds the “best” f on the training data
- ▶ Same idea applies for unsupervised problem (minimizing reconstruction error)

Empirical risk minimization

Train by minimizing the train error

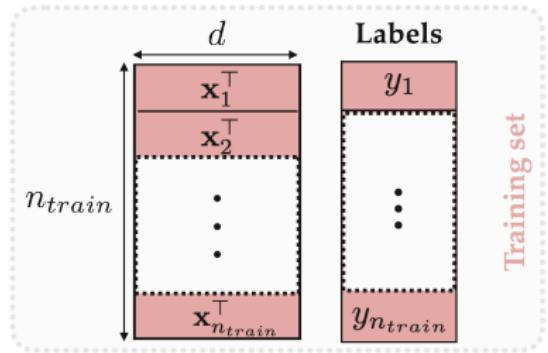
To find f the idea is to **minimize the averaged error** on the training samples:

$$\min_f \frac{1}{n} \sum_{i=1}^n \ell(y_i, f(\mathbf{x}_i)) \quad (\text{ERM})$$

- ▶ It is called **empirical risk minimization (ERM)**
- ▶ Given the loss, finds the “best” f on the training data
- ▶ Same idea applies for unsupervised problem (minimizing reconstruction error)

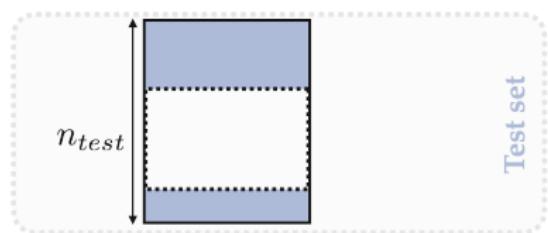
Once solved how do I know if my model is good ?

Train/validation/test

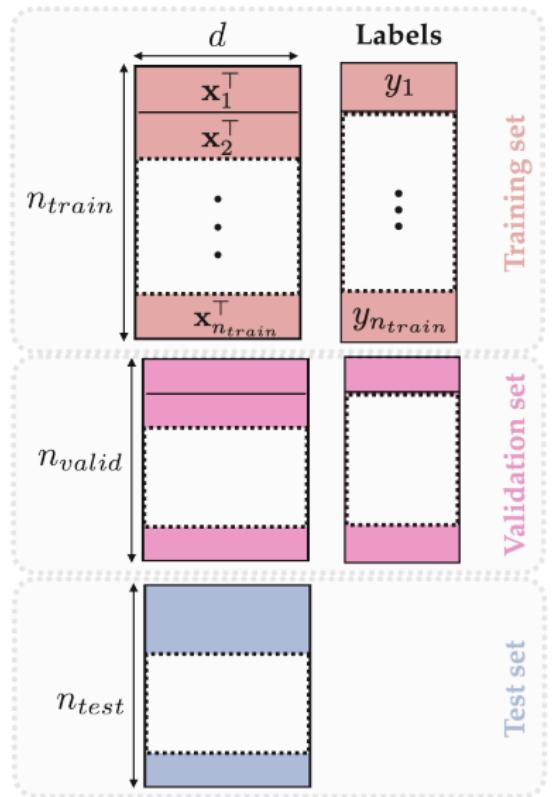


Train, validation and test data

- ▶ Data are divided in two sets: a train and a test set
- ▶ Learning a ML model = on the train set where you know the output of the model
- ▶ The test set is **never available**, used only to deploy the model



Train/validation/test

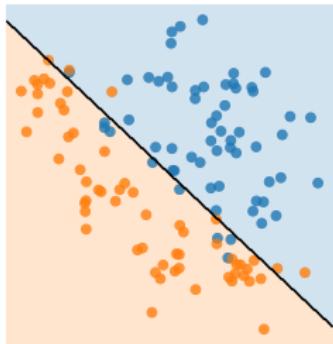


Train, validation and test data

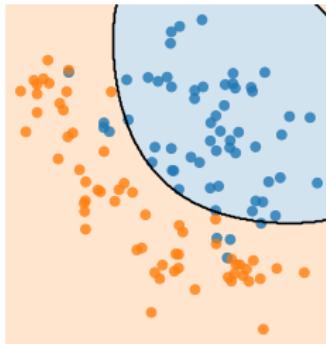
- ▶ Data are divided in two sets: a train and a test set
- ▶ Learning a ML model = on the train set where you know the output of the model
- ▶ The test set is **never available**, used only to deploy the model
- ▶ A validation set is often used to **calibrate the model** (choose hyperparameters), and **evaluate the generalization properties** of your model

Underfitting and overfitting

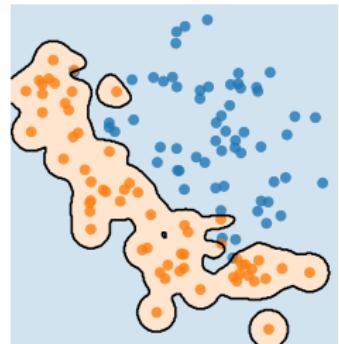
Acc. 0.89/0.89 train/test



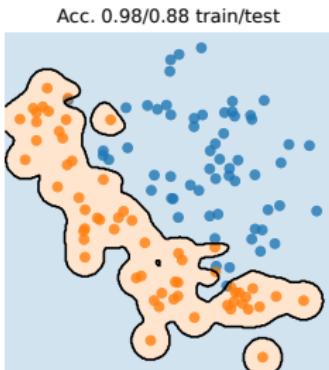
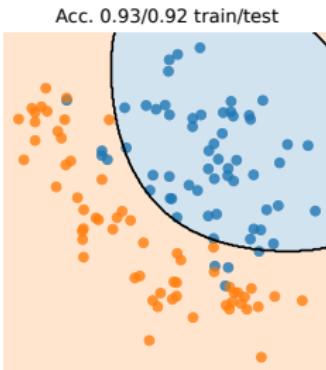
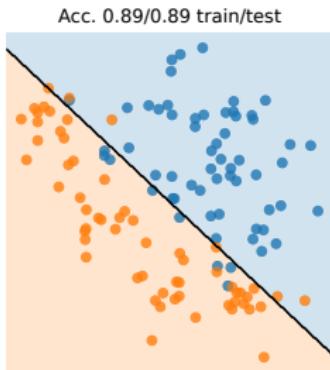
Acc. 0.93/0.92 train/test



Acc. 0.98/0.88 train/test

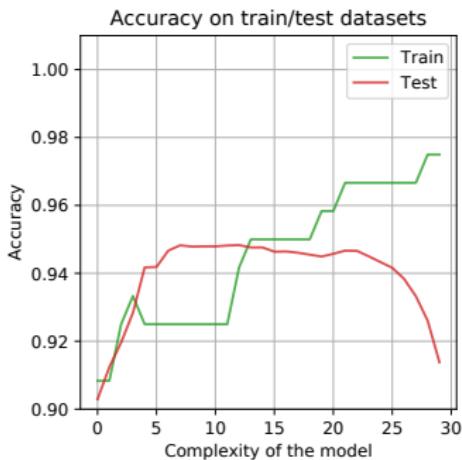


Underfitting and overfitting



Complexity of a model

- ▶ Under-fitting when the model is too simple.
- ▶ Over-fitting occurs when the model is too complex
- ▶ Training data performance is not a good proxy for testing performance.
- ▶ We want to predict well on new data!
- ▶ Parameter and model validation.



Empirical risk minimization

Train by minimizing the train error

To find f the idea is to **minimize the averaged error** on the training samples:

$$\min_f \frac{1}{n} \sum_{i=1}^n \ell(y_i, f(\mathbf{x}_i)) + \lambda \text{Reg}(f) \quad (\text{ERM})$$

- ▶ It is called **empirical risk minimization (ERM)**
- ▶ Given the loss, finds the “best” f on the training data
- ▶ Teacher/student analogy
- ▶ Same idea applies for unsupervised problem (minimizing reconstruction error)

... but we want generalization !

- ▶ We want f to be good outside the training samples
- ▶ Add regularization ! (limit the complexity of f)

Plan

What is machine learning ?

Data in machine learning

From training data to prediction

Loss functions

Empirical risk minimization

Train/validation/test

Model selection and validation

Split your dataset !

The problems with structured data

Motivating examples

A primer on graph theory

Why “classical ML” struggles with structured data

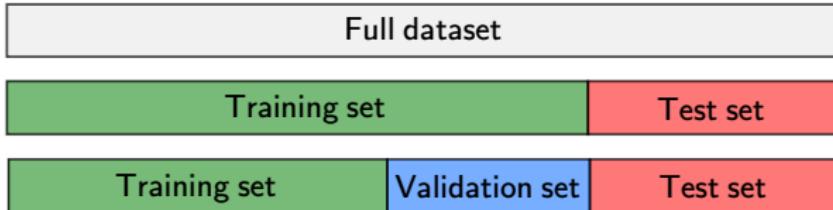
Model selection and validation

Bias-complexity tradeoff

$$\text{excess risk} = \text{estimation error} + \text{approximation error}$$

Conclusion: we have to select a model that is not too complex but not too simple !

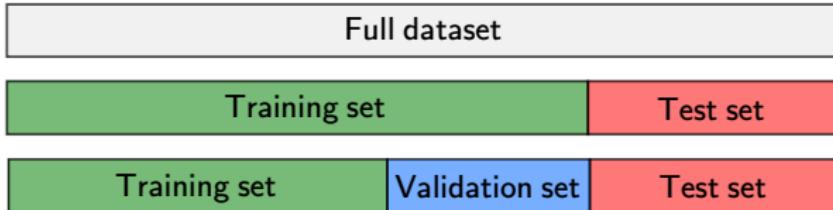
Splitting the data



Principle of Hold-Out cross-validation

- ▶ Split the training data in a training and validation sets (non overlapping).
- ▶ Train different models (different methods and/or parameters) on the train data.
- ▶ Evaluate performance on the validation data and select the method/parameters with best performance.

Splitting the data



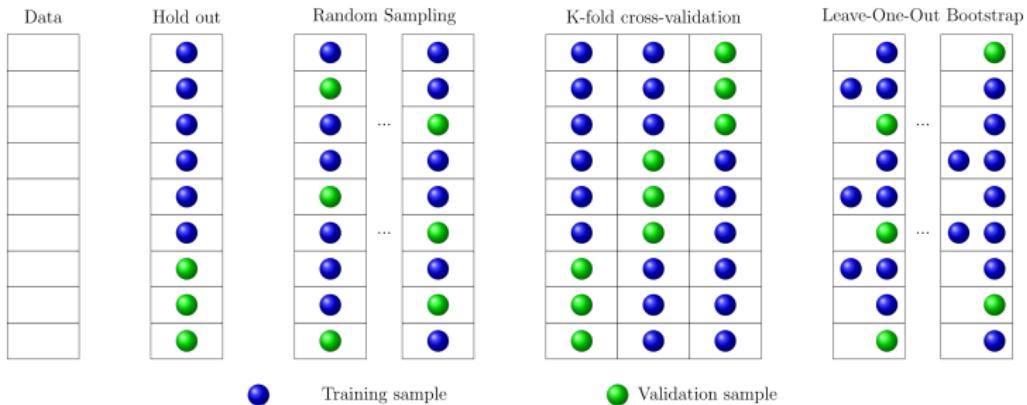
Principle of Hold-Out cross-validation

- ▶ Split the training data in a training and validation sets (non overlapping).
- ▶ Train different models (different methods and/or parameters) on the train data.
- ▶ Evaluate performance on the validation data and select the method/parameters with best performance.

Final estimator

- ▶ After selecting the optimal parameters, one should retrain the estimator on the whole training dataset using the optimal method/parameters.

Different ways to split the data



Data splitting for cross-validation Arlot and Celisse 2010

- ▶ The training data is split in non-overlapping training/validation sets.
- ▶ **Hold-Out** uses a unique split and computes the performance on the validation set.
- ▶ More robust cross-validation approaches actually investigate several splits of the data and compute the average performance:
 - ▶ **K-fold** (split in K sets and use one split as test for all k)
 - ▶ Random sampling (aka **Shuffle split**) draws several random splittings.
- ▶ Scikit-learn implementation : `sklearn.model_selection.cross_valide`

Plan

What is machine learning ?

Data in machine learning

From training data to prediction

Loss functions

Empirical risk minimization

Train/validation/test

Model selection and validation

Split your dataset !

The problems with structured data

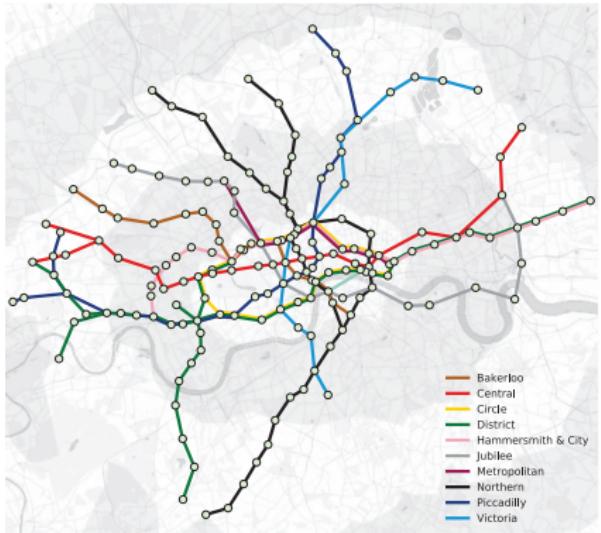
Motivating examples

A primer on graph theory

Why “classical ML” struggles with structured data

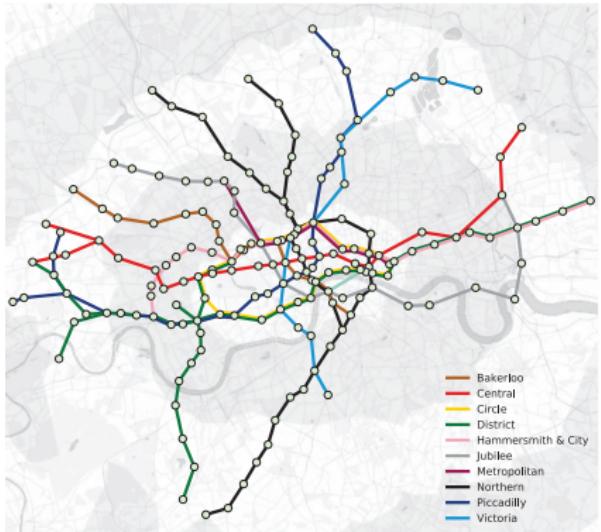
Motivating examples

- ▶ Traffic forecasting (e.g. ETA estimation): GNN for Google Maps Derrow-Pinion et al. 2021.



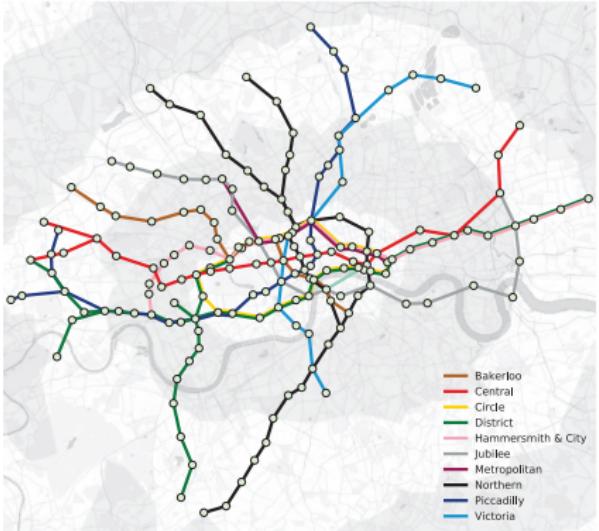
Motivating examples

- ▶ Traffic forecasting (e.g. ETA estimation): GNN for Google Maps [Derrow-Pinion et al. 2021](#).
- ▶ Chemistry and Drug Design: space of chemically synthesisable molecules is very large (estimated around 10^{60}).
- ▶ Drug Repositioning: action of drugs and their interactions → graph [Barabási, Gulbahce, and Loscalzo 2011](#).

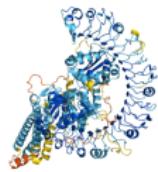


Motivating examples

- ▶ Traffic forecasting (e.g. ETA estimation): GNN for Google Maps [Derrow-Pinion et al. 2021](#).
- ▶ Chemistry and Drug Design: space of chemically synthesisable molecules is very large (estimated around 10^{60}).
- ▶ Drug Repositioning: action of drugs and their interactions → graph [Barabási, Gulbahce, and Loscalzo 2011](#).
- ▶ Protein biology [Jumper et al. 2021](#).

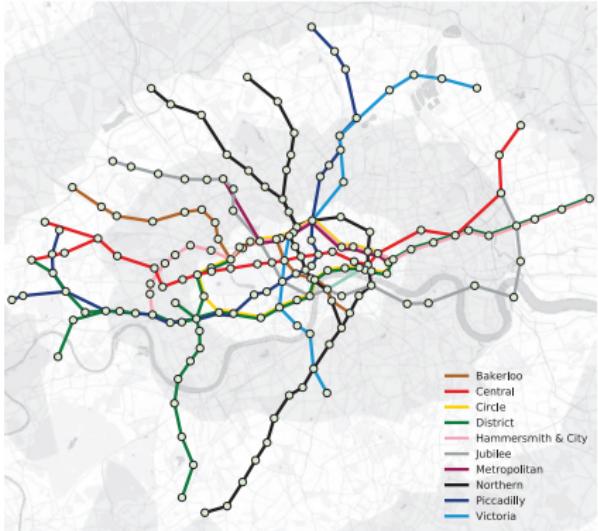


AAATGCG.... →



Motivating examples

- ▶ Traffic forecasting (e.g. ETA estimation): GNN for Google Maps [Derrow-Pinion et al. 2021](#).
- ▶ Chemistry and Drug Design: space of chemically synthesisable molecules is very large (estimated around 10^{60}).
- ▶ Drug Repositioning: action of drugs and their interactions → graph [Barabási, Gulbahce, and Loscalzo 2011](#).
- ▶ Protein biology [Jumper et al. 2021](#).
- ▶ Recommender Systems and Social Networks.
- ▶ Healthcare.

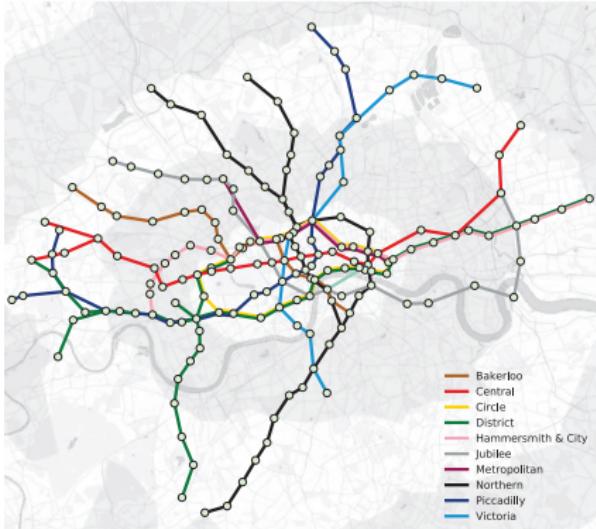


AAATGCG.... →

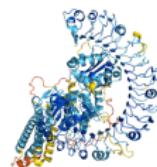


Motivating examples

- ▶ Traffic forecasting (e.g. ETA estimation): GNN for Google Maps [Derrow-Pinion et al. 2021](#).
- ▶ Chemistry and Drug Design: space of chemically synthesisable molecules is very large (estimated around 10^{60}).
- ▶ Drug Repositioning: action of drugs and their interactions → graph [Barabási, Gulbahce, and Loscalzo 2011](#).
- ▶ Protein biology [Jumper et al. 2021](#).
- ▶ Recommender Systems and Social Networks.
- ▶ Healthcare.
- ▶ and more...



AAATGCG.... →



What is a graph ?

Definition

A graph $G = (V, E)$ is defined as a set of **vertices** V , which are connected by a set of **edges** $E \subset V \times V$.

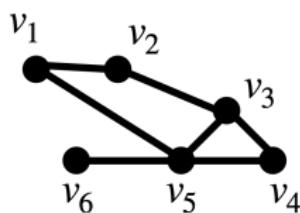
What is a graph ?

Definition

A graph $G = (V, E)$ is defined as a set of **vertices** V , which are connected by a set of **edges** $E \subset V \times V$.

- ▶ Example of **undirected** graph

$$G = (V, E)$$



Adjacency matrix

$$\mathbf{A} = \begin{pmatrix} v_1 & v_2 & v_3 & v_4 & v_5 & v_6 \\ 0 & 1 & 0 & 0 & 1 & 0 \\ 1 & 0 & 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 1 & 1 & 0 \\ 0 & 0 & 1 & 0 & 1 & 0 \\ 1 & 0 & 1 & 1 & 0 & 1 \\ 0 & 0 & 0 & 0 & 1 & 0 \end{pmatrix} = \mathbf{A}^T$$

Adjacency matrix

The adjacency $\mathbf{A} \in \mathbb{R}^{|V| \times |V|}$ is defined as

$$[\mathbf{A}]_{ij} = \begin{cases} 1 & \text{if } (v_i, v_j) \in E \text{ (often noted as } v_i \sim v_j\text{)} \\ 0 & \text{otherwise} \end{cases}$$

What is a graph ?

Definition

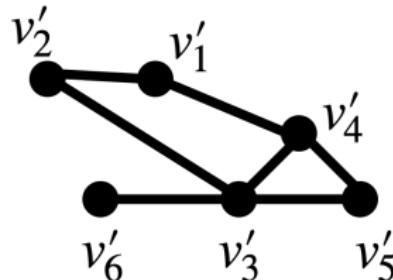
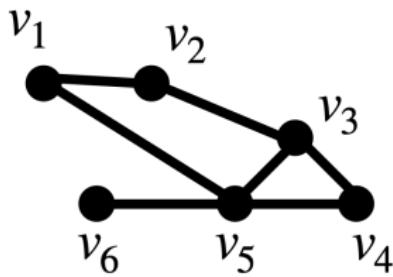
A graph $G = (V, E)$ is defined as a set of **vertices** V , which are connected by a set of **edges** $E \subset V \times V$.

Isomorphic graphs

The definition depends on the ordering of the nodes.

$$G = (V, E)$$

$$G' = (V', E')$$



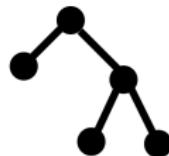
$$\mathbf{A} = \mathbf{P}^T \mathbf{A}' \mathbf{P}$$

What is a graph ?

Definition

A graph $G = (V, E)$ is defined as a set of **vertices** V , which are connected by a set of **edges** $E \subset V \times V$.

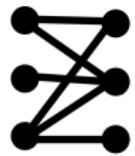
Some special structures



Tree



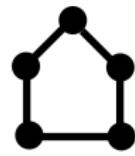
Complete graph



Bipartite graph



Star graph

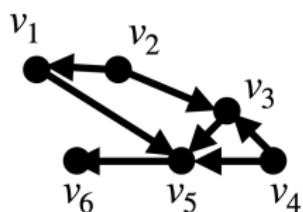


Circular graph

What is a graph ?

- ▶ Example of **directed** graph

$$G = (V, E)$$

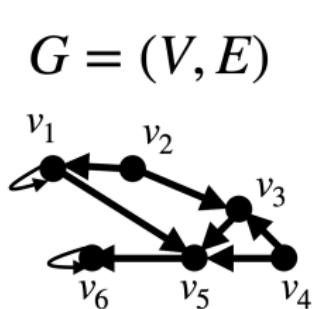


Adjacency matrix

$$\mathbf{A} = \begin{pmatrix} v_1 & v_2 & v_3 & v_4 & v_5 & v_6 \\ 0 & 0 & 0 & 0 & 1 & 0 \\ 1 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 1 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{pmatrix} \neq \mathbf{A}^T$$

What is a graph ?

- ▶ Example of **directed** graph with **self-loops**.



Adjacency matrix

$$A = \begin{pmatrix} v_1 & v_2 & v_3 & v_4 & v_5 & v_6 \\ 1 & 0 & 0 & 0 & 1 & 0 \\ 1 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 1 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 & 0 & 1 \end{pmatrix} \neq A^T$$

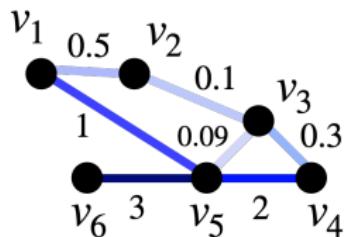
What is a graph ?

Weighted graph

A weighted graph $G = (V, E)$ associates non-negative weights to each edge.

- ▶ Example of **weighted** graph

$$G = (V, E)$$



Weight matrix

$$\mathbf{W} = \begin{pmatrix} 0 & 0.5 & 0 & 0 & 1 & 0 \\ 0.5 & 0 & 0.1 & 0 & 0 & 0 \\ 0 & 0.1 & 0 & 0.3 & 0.09 & 0 \\ 0 & 0 & 0.3 & 0 & 2 & 0 \\ 1 & 0 & 0.09 & 2 & 0 & 3 \\ 0 & 0 & 0 & 0 & 3 & 0 \end{pmatrix}$$

What is a graph ?

Degree of a node

The degree of a node v_i is

$$d_i = |\{v \in V : v \sim v_i\}| = \sum_{j=1}^{|V|} A_{ij}$$

The degree matrix is $\mathbf{D} = \text{diag}(d_1, \dots, d_{|V|})$.

What is a graph ?

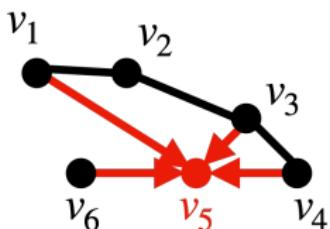
Degree of a node

The degree of a node v_i is

$$d_i = |\{v \in V : v \sim v_i\}| = \sum_{j=1}^{|V|} A_{ij}$$

The degree matrix is $\mathbf{D} = \text{diag}(d_1, \dots, d_{|V|})$.

$$G = (V, E)$$



Adjacency matrix Degree

$$\mathbf{A} = \begin{pmatrix} v_1 & v_2 & v_3 & v_4 & v_5 & v_6 \\ 0 & 1 & 0 & 0 & 1 & 0 \\ 1 & 0 & 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 1 & 1 & 0 \\ 0 & 0 & 1 & 0 & 1 & 0 \\ \textcolor{red}{1} & \textcolor{red}{0} & \textcolor{red}{1} & \textcolor{red}{1} & \textcolor{red}{0} & \textcolor{red}{1} \\ 0 & 0 & 0 & 0 & 1 & 0 \end{pmatrix}$$
$$\begin{pmatrix} 2 \\ 2 \\ 3 \\ 2 \\ 4 \\ 1 \end{pmatrix} = \mathbf{A}\mathbf{1}$$

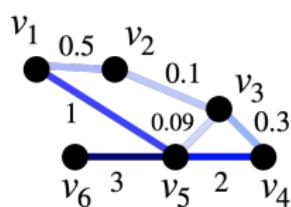
What is a graph ?

Degree of a node

The degree of a node v_i in a weighted graph is

$$d_i = \sum_{j=1}^{|V|} W_{ij}$$

$$G = (V, E)$$



	Weight matrix						Degree
$W =$	0	0.5	0	0	1	0	$\begin{pmatrix} 1.5 \\ \end{pmatrix}$
	0.5	0	0.1	0	0	0	$\begin{pmatrix} 0.6 \\ \end{pmatrix}$
	0	0.1	0	0.3	0.09	0	$\begin{pmatrix} 0.49 \\ \end{pmatrix}$
	0	0	0.3	0	2	0	$\begin{pmatrix} 2.3 \\ \end{pmatrix}$
	1	0	0.09	2	0	3	$\begin{pmatrix} 4.09 \\ \end{pmatrix}$
	0	0	0	0	3	0	$\begin{pmatrix} 3 \\ \end{pmatrix}$

What is a graph ?

Laplacian matrix

The Laplacian matrix of a undirected graph is defined as

$$\mathbf{L} = \mathbf{D} - \mathbf{W} \text{ where } \mathbf{D} \text{ is the degree matrix}$$

Properties

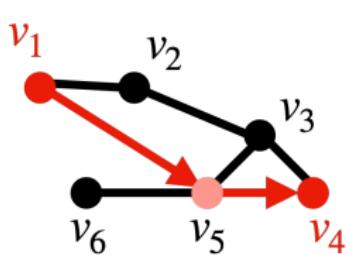
On the board

What is a graph ?

Shortest-path matrix

The shortest-path between $v, v' \in V$ is the path that connects v, v' such that the sum of the weights of its constituent edges is minimized.

$$G = (V, E)$$



Shortest-path matrix

$$S = \begin{pmatrix} 0 & 1 & 2 & 2 & 1 & 2 \\ 1 & 0 & 1 & 2 & 2 & 3 \\ 2 & 1 & 0 & 1 & 1 & 2 \\ 2 & 2 & 1 & 0 & 1 & 2 \\ 1 & 2 & 1 & 1 & 0 & 1 \\ 2 & 3 & 2 & 2 & 1 & 0 \end{pmatrix}$$

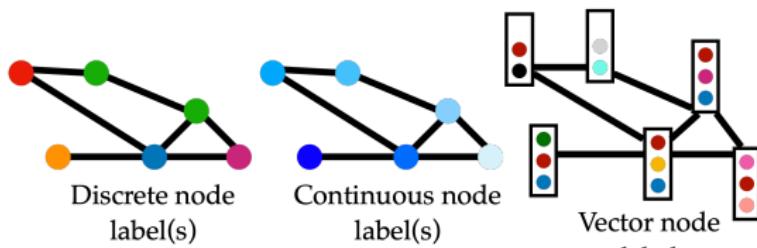
- ▶ Dijkstra's algorithm computes all the shortest paths from a single node in $O(|E| + |V| \log(|V|))$.
- ▶ All-pairs shortest paths with Floyd–Warshall algorithm in $O(|V|^3)$.

What is a graph ?

Attributed graphs

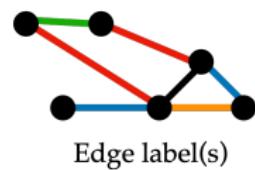
Most graphs encountered in ML also have **attributes**.

Attributed graphs



$$\ell_G : V \rightarrow S \text{ (} \subset \mathbb{R}^N \text{)}$$

Edge attributed graphs



Edge label(s)

$$\ell_E : E \rightarrow S_E$$

ML vs structured data

Problems

ML vs structured data

Problems

- ▶ Can we encode a graph G as a vector $\in \mathbb{R}^d$ to use standard ML methods ?

ML vs structured data

Problems

- ▶ Can we encode a graph G as a vector $\in \mathbb{R}^d$ to use standard ML methods ?
- ▶ Can we build ML methods with the raw representation of G ? How to adapt ML methods that work on vectors ?

ML vs structured data

Problems

- ▶ Can we encode a graph G as a vector $\in \mathbb{R}^d$ to use standard ML methods ?
- ▶ Can we build ML methods with the raw representation of G ? How to adapt ML methods that work on vectors ?
- ▶ How can we handle the combinatoric nature of graphs ?

ML vs structured data

Problems

- ▶ Can we encode a graph G as a vector $\in \mathbb{R}^d$ to use standard ML methods ?
- ▶ Can we build ML methods with the raw representation of G ? How to adapt ML methods that work on vectors ?
- ▶ How can we handle the combinatoric nature of graphs ?
- ▶ ML outputs should be permutation invariant ? equivariant ?

ML vs structured data

Problems

- ▶ Can we encode a graph G as a vector $\in \mathbb{R}^d$ to use standard ML methods ?
- ▶ Can we build ML methods with the raw representation of G ? How to adapt ML methods that work on vectors ?
- ▶ How can we handle the combinatoric nature of graphs ?
- ▶ ML outputs should be permutation invariant ? equivariant ?
- ▶ When data = vectors in one graph how can we take into account the structure of the graph ?

ML vs structured data

Problems

- ▶ Can we encode a graph G as a vector $\in \mathbb{R}^d$ to use standard ML methods ?
- ▶ Can we build ML methods with the raw representation of G ? How to adapt ML methods that work on vectors ?
- ▶ How can we handle the combinatoric nature of graphs ?
- ▶ ML outputs should be permutation invariant ? equivariant ?
- ▶ When data = vectors in one graph how can we take into account the structure of the graph ?
- ▶ When data = vectors can we find an (interesting) graph that represent these data ?

References I

-  Arlot, Sylvain and Alain Celisse (2010). "A survey of cross-validation procedures for model selection". In: *Statistics Surveys* 4.
-  Bach, Francis (2022). *Learning Theory from First Principles*.
-  Barabási, Albert-László, Natali Gulbahce, and Joseph Loscalzo (2011). "Network medicine: a network-based approach to human disease". In: *Nature reviews genetics* 12.1, pp. 56–68.
-  Derrow-Pinion, Austin et al. (2021). "Eta prediction with graph neural networks in google maps". In: *Proceedings of the 30th ACM International Conference on Information & Knowledge Management*, pp. 3767–3776.
-  Hastie, Trevor, Robert Tibshirani, and Jerome Friedman (2001). *The Elements of Statistical Learning*. Springer New York Inc.
-  Jumper, John et al. (2021). "Highly accurate protein structure prediction with AlphaFold". In: *Nature* 596.7873, pp. 583–589.
-  Shalev-Shwartz, Shai and Shai Ben-David (2014). *Understanding Machine Learning - From Theory to Algorithms*. Cambridge University Press.